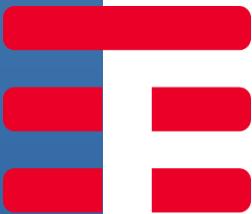




GRUPPO TELECOM ITALIA
Big Data Transformation

Big Data Transformation

TIM Academy



GRUPPO TIM

Big Data Transformation

Apache Spark Fundamentals & Query Fundamentals

 **TIM** Academy

The TIM Academy logo features a red stylized 'T' icon followed by the word 'Academy' in a white, sans-serif font.

GRUPPO TIM

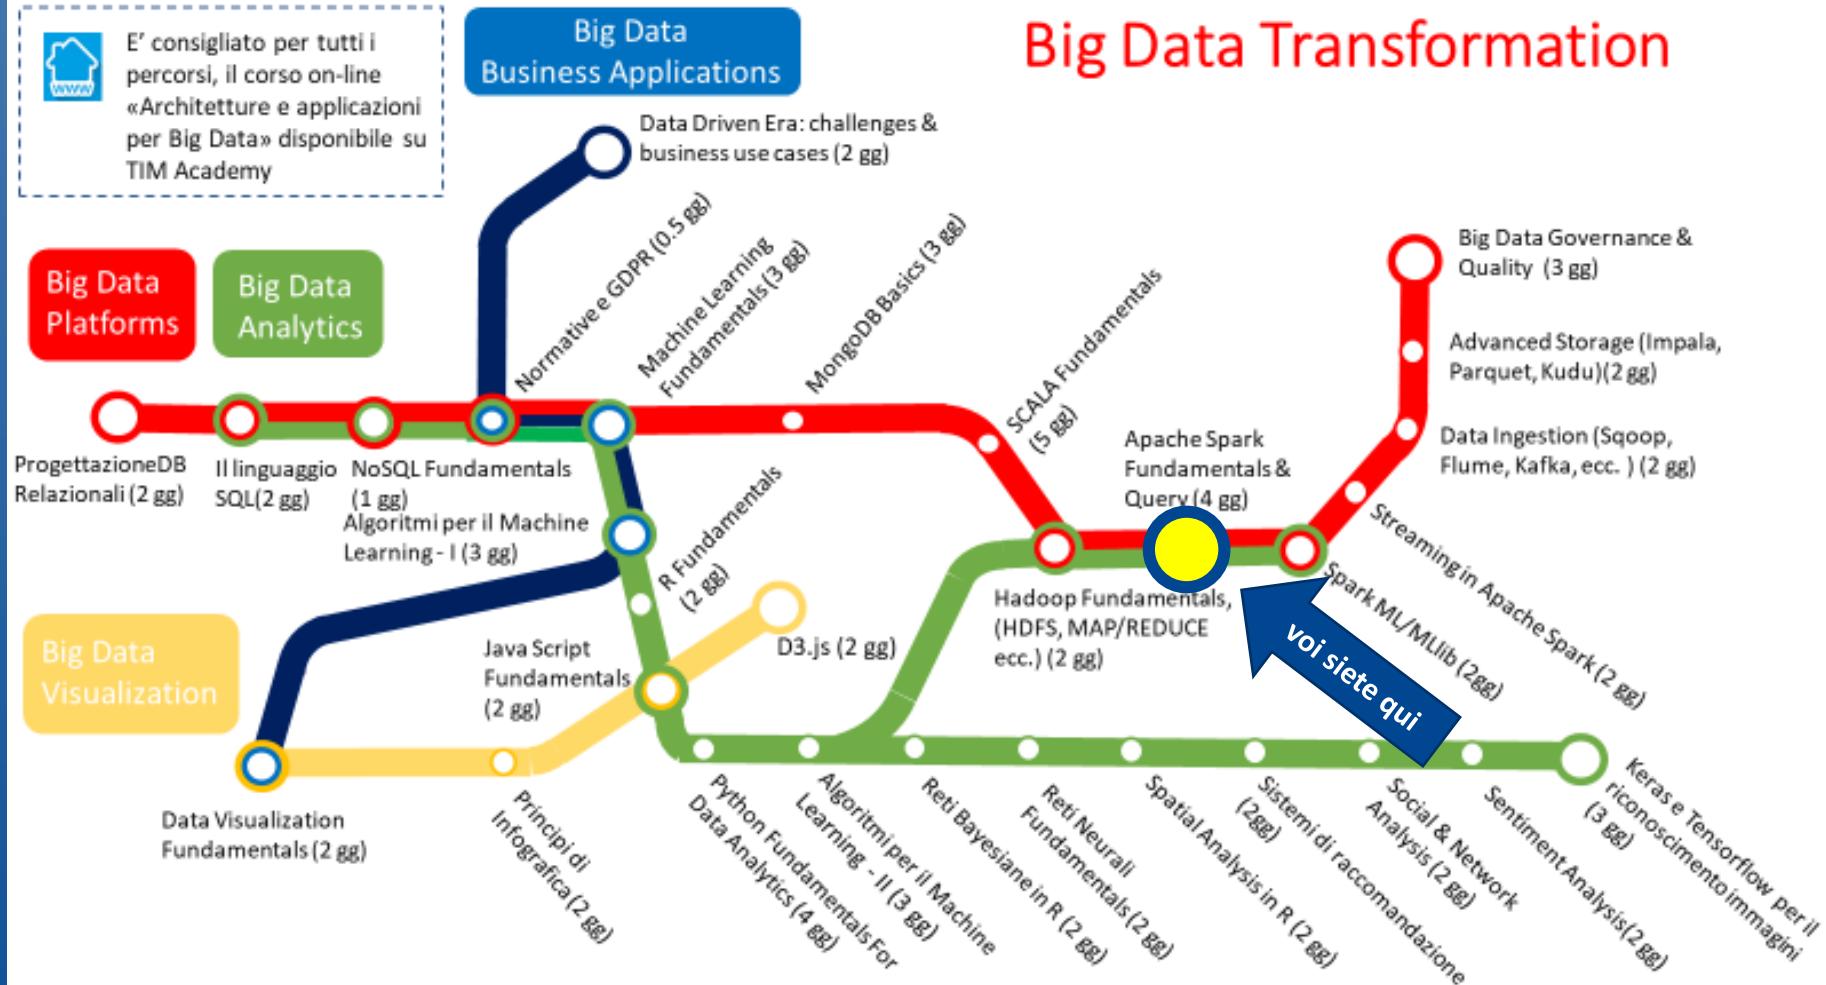
Big Data Transformation

Apache Spark Fundamentals & Query

 **TIM Academy**



E' consigliato per tutti i percorsi, il corso on-line «Architetture e applicazioni per Big Data» disponibile su TIM Academy

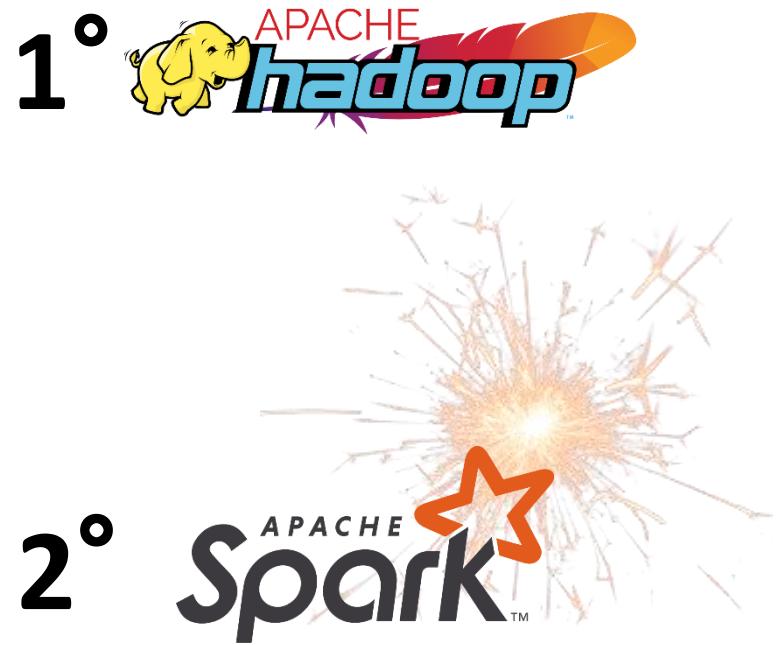


Apache Spark

Uno dei tool più usati nel campo dei **BIG DATA**



<https://www.whizlabs.com/blog/big-data-tools/>



Prerequisiti

- Conoscenza del linguaggio di programmazione Scala



- Aver seguito il corso Hadoop Fundamentals



Agenda del corso

| | 1° giorno | 2° giorno | 3° giorno | 4° giorno | |
|------------|-------------------------------------|-------------------------------------|-------------------------------------|-------------------------------------|--------------------------------------------------|
| mattina | 9:00→ 9:30 → →13:00 14:00→ | 9:00→ →13:00 14:00→ →17:30 | 9:00→ →13:00 14:00→ →17:30 | 9:00→ →13:00 14:00→ →17:30 | lezione laboratorio lezione laboratorio |
| pomeriggio | →17:30 | | | | |

Agenda 1/8

— Introduzione

Presentazione del corso

Spark

Cos'è? A cosa serve? Chi lo usa? Punti di forza e debolezze

Contesto e storia

Come è nato? Chi l'ha fatto? Quale problema affronta?

Com'è fatto?

Una prima presentazione dell'architettura

E in pratica?

Come si usa? Come lo useremo noi in laboratorio?



Agenda 2/8

— Le fondamenta

La programmazione funzionale

Un paradigma utile per il calcolo parallelo

Strutture dati

RDDs, Datasets e DataFrames

Operazioni sui dati

Transformazioni e Azioni; Valutazioni «pigre»; persistenza (cache)

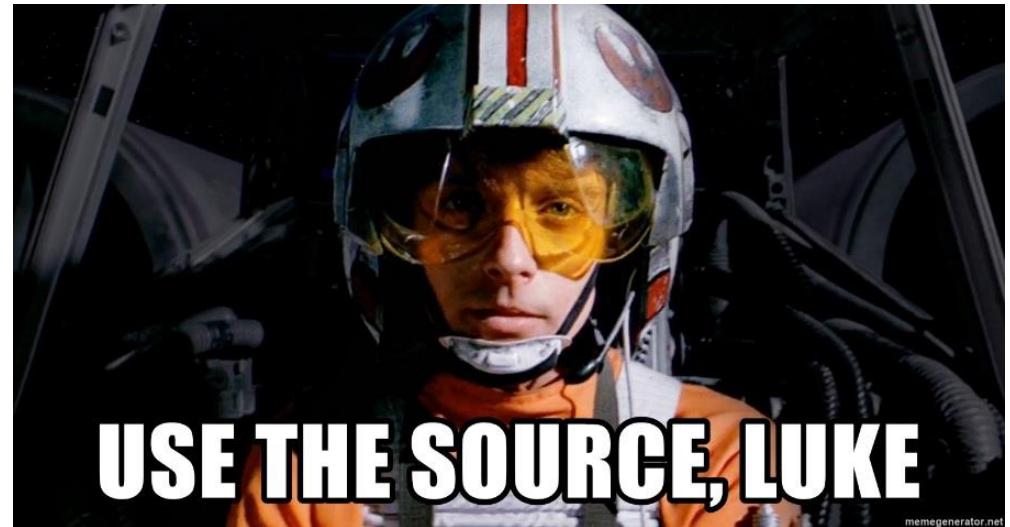
Punti di riferimento

Documentazione; Cloudera, Hortonworks; Databricks;

Codice sorgente su github

Esercizi di laboratorio

Mappe, filtri e qualche azione



Agenda 3/8

— API Strutturata

DataFrames e Dataset in profondità

Data Source

Date, Stringhe ed Espressioni regolari

Union, Join e SQL

```
root
|-- InvoiceNo: string (nullable = true)
|-- StockCode: string (nullable = true)
|-- Description: string (nullable = true)
|-- Quantity: integer (nullable = true)
|-- InvoiceDate: timestamp (nullable = true)
|-- UnitPrice: double (nullable = true)
|-- CustomerID: double (nullable = true)
|-- Country: string (nullable = true)
```

| InvoiceNo | StockCode | Description | Quantity | InvoiceDate | UnitPrice | CustomerID | Country |
|-----------|-----------|----------------------|----------|---------------------|-----------|------------|---------|
| 536365 | 85123A | WHITE HANGING HEA... | 6 | 2010-12-01 08:26:00 | ... | | |
| 536365 | 71053 | WHITE METAL LANTERN | 6 | 2010-12-01 08:26:00 | ... | | |
| ... | | | | | | | |
| 536367 | 21755 | LOVE BUILDING BLO... | 3 | 2010-12-01 08:34:00 | ... | | |
| 536367 | 21777 | RECIPE BOX WITH M... | 4 | 2010-12-01 08:34:00 | ... | | |

Agenda 4/8

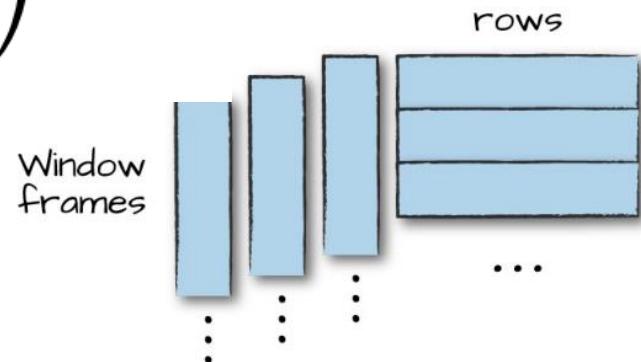
— Aggregazioni

Tipi diversi di aggregazione

Window

Aggregazioni su tipi complessi

$$\left(\sum_{i=1}^n p_i x_i^2 \right) - \mu^2$$



Agenda 5/8

RDD

La prima API e quella fondamentale

Tipi diversi di RDD. RDD di coppie chiave-valore

Partizioni

Variabili condivise: accumulatori e variabili broadcast

RDD

Agenda 6/8

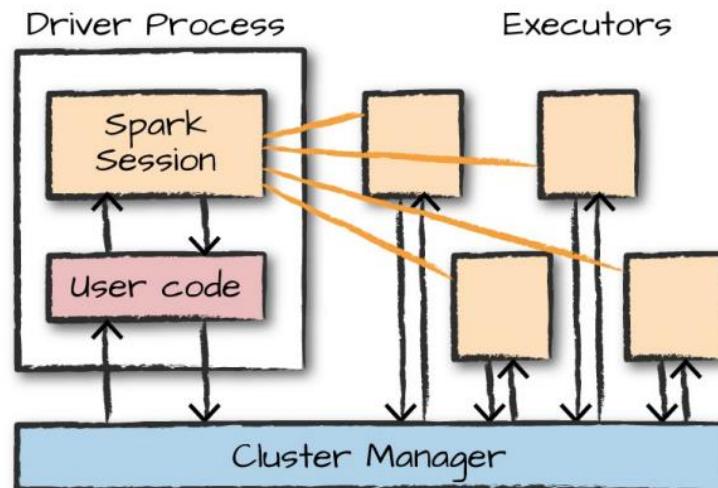
Il cluster

Struttura di un'applicazione Spark

Modi di esecuzione

Ciclo di vita di un'applicazione Spark

Sessione, programma e configurazione



Agenda 7/8

— GraphFrames

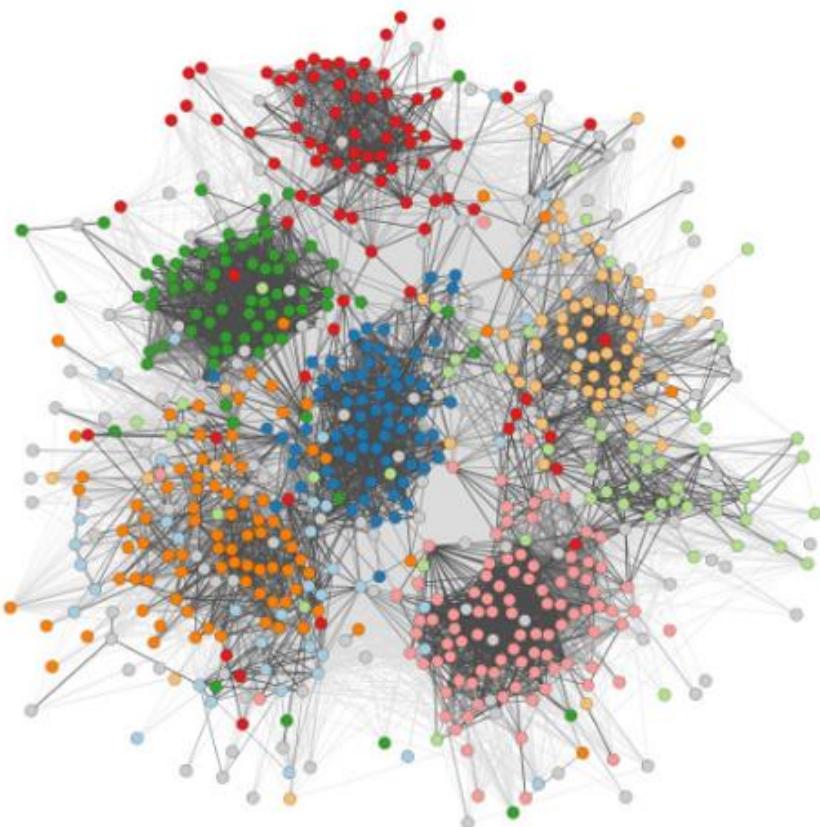
Packages

Spark streaming, Spark ML

GraphX e GraphFrames

Cos'è un grafo? Vertici e Archi

Algoritmi principali



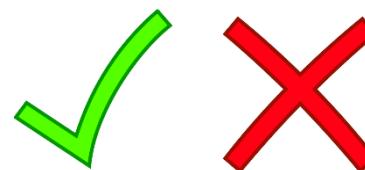
Agenda 8/8

— Conclusioni

«Do» e «don't» in Spark

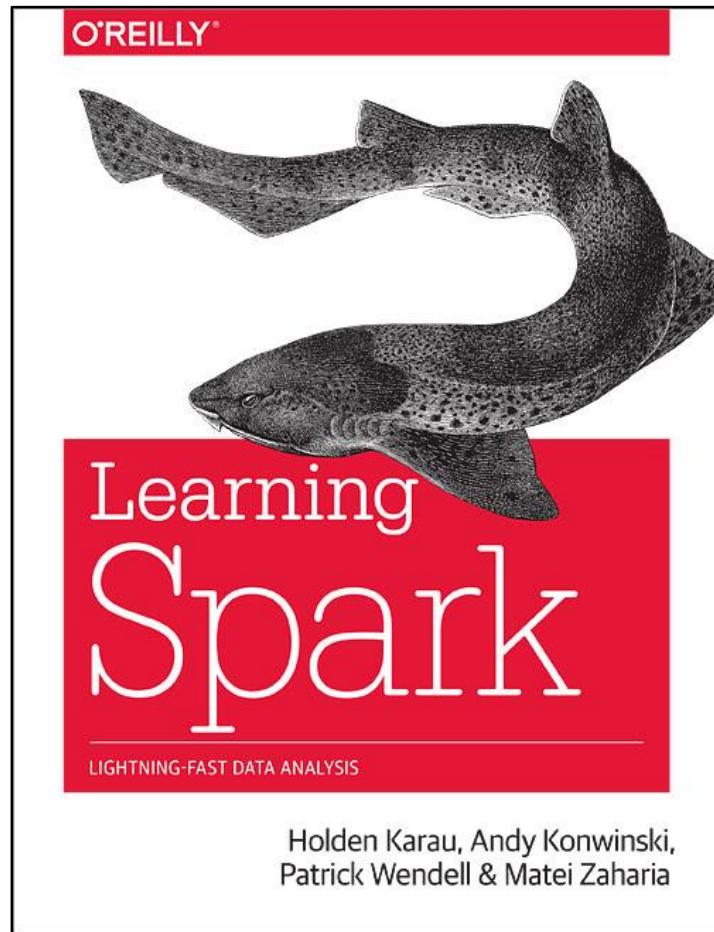
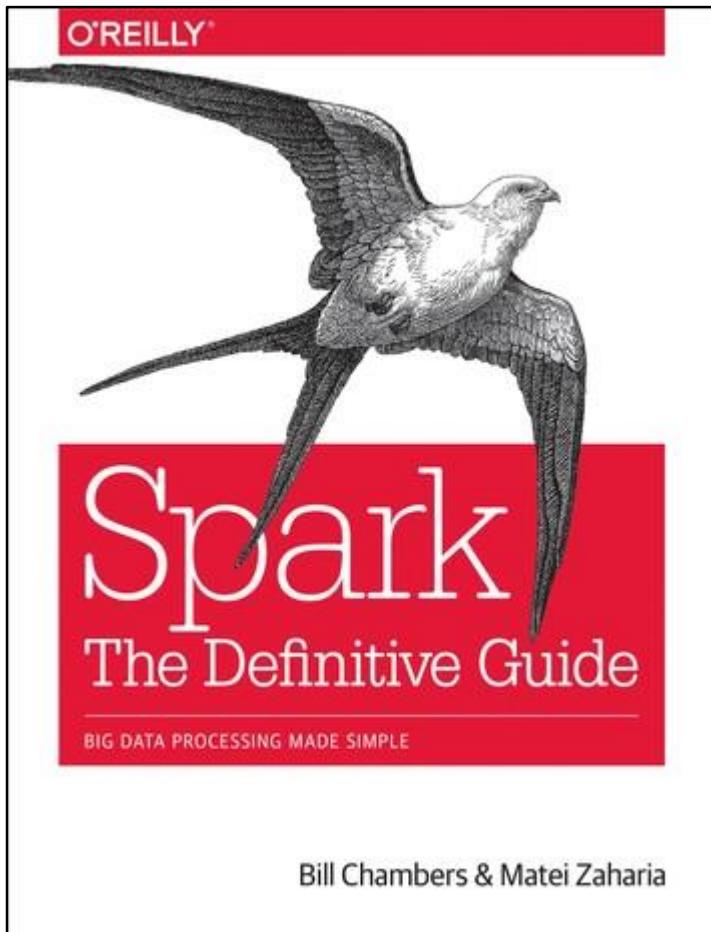
Recapitolazione e domande

Compilazione del questionario





Riferimenti bibliografici



Cos'è Apache Spark?

Cos'è Apache Spark?



WIKIPEDIA
The Free Encyclopedia

“Apache Spark is an **open-source distributed general-purpose cluster-computing framework**”.

- **Open-source**

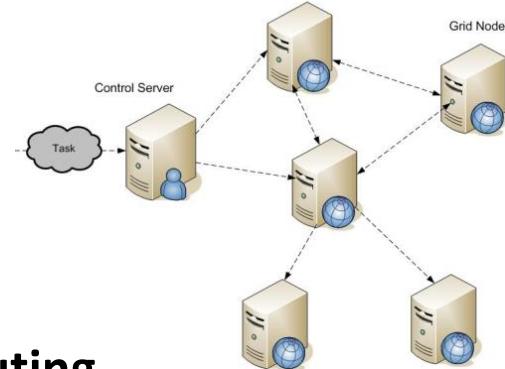


Open Source
Initiative

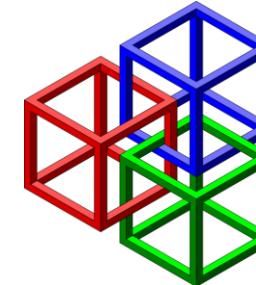
- **General-purpose**



- **Distributed**
- **Cluster-computing**



- **Framework**



Computing... ad esempio?

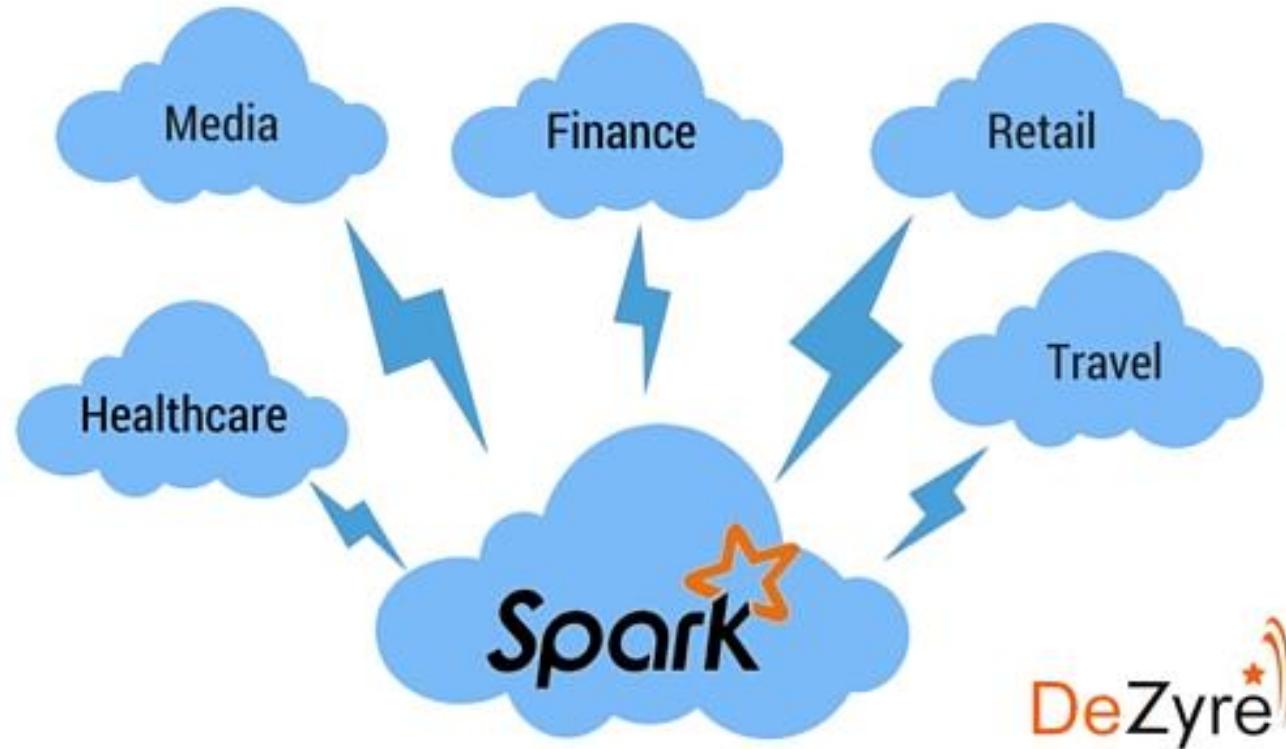
Vari utilizzi di Apache Spark

- Streaming
 - Streaming ETL (Extract, Transform, Load)
 - Data enrichment
 - Event detection (ad es. per individuare transizioni fraudolente, cambi potenzialmente pericolosi dei parametri vitali, ecc.)
 - Complex session analysis
- Machine learning
 - Marketing purposes and sentiment analysis
 - Network security
- Interactive analysis

<https://www.qubole.com/blog/apache-spark-use-cases/>

Computing... ad esempio?

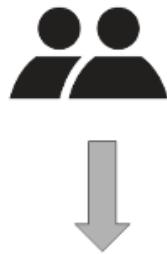
Vari utilizzi di Apache Spark



<https://www.dezyre.com/article/top-5-apache-spark-use-cases/271>

Computing... ad esempio?

Test A/B



Welcome to our website

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

[Learn more](#)

Click rate: **52 %**



Welcome to our website

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

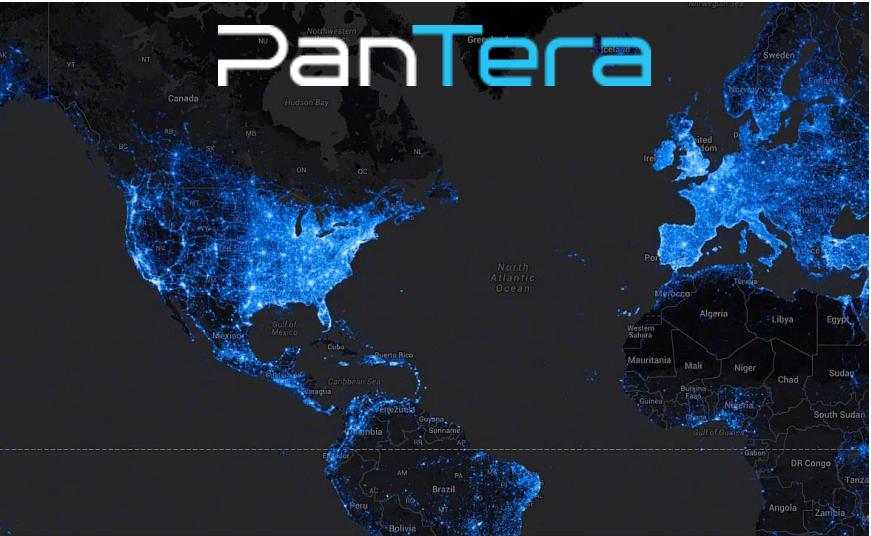
[→ Learn more](#)

72 %

Qualche esempio concreto



Jet Propulsion Laboratory
California Institute of Technology

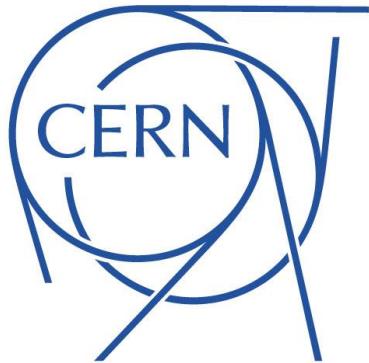


YAHOO!

Rif: <https://spark.apache.org/powerd-by.html>

Qualche esempio concreto

NETFLIX



ebaytm

Uber

ING The ING logo features the word "ING" in a dark blue serif font next to a golden orange lion rampant.

DBS

GE Aviation

BLACKHAWK
N E T W O R K

NBC **UNIVERSAL**

...

Fonte: <https://databricks.com/blog/2018/09/05/a-guide-apache-spark-use-cases-streaming-and-research-talks-at-spark-ai-summit-europe.html>
<https://www.infoworld.com/article/3031690/analytics/why-you-should-use-spark-for-machine-learning.amp.html>

Qualche esempio concreto



“[...] eBay is migrating its 30 PB MPP database to Apache Spark [...]”

Massively Parallel Processing.
1 PetaByte = 1.000.000 GB

[2018]

<https://databricks.com/session/experience-of-optimizing-spark-sql-when-migrating-from-mpp-database>



monitor or detect or predict preventive maintenance

[2018]

“[...] Analytic tools such as SQL Server and MATLAB were used until recently, when GE’s data was moved to an Apache Spark environment. Consequently, our advanced analytics are now being migrated to Spark, [...]”

<https://databricks.com/session/ge-aviation-spark-application-experience-porting-analytics-into-pyspark-ml-pipelines>



[2015?]

“In the system presented, we run Spark to run the data analytics pipeline for anomaly detection. [...] Our machine learning pipeline uses Spark **decision tree ensembles** and **k-means clustering**.”

<https://spark-summit.org/eu-2015/events/real-time-anomaly-detection-with-spark-ml-and akka/>

Fonte: <https://databricks.com/blog/2018/09/05/a-guide-apache-spark-use-cases-streaming-and-research-talks-at-spark-ai-summit-europe.html>
<https://www.infoworld.com/article/3031690/analytics/why-you-should-use-spark-for-machine-learning.amp.html>

Qualche esempio concreto



Dal 2015 usa *Hadoop + Kafka + Spark* (ecc.)
100+ Petabytes con un minuto di latenza
10 000 job **Spark** / giorno

Cfr: <https://eng.uber.com/uber-big-data-platform/>



175 milioni di utenti mensilmente attivi
Test A/B
Kafka + Pinball + Spark Streaming + HBase & MemSQL + Presto

Cfr: https://medium.com/@Pinterest_Engineering/building-a-new-experiment-pipeline-with-spark-b04e1a7d639a



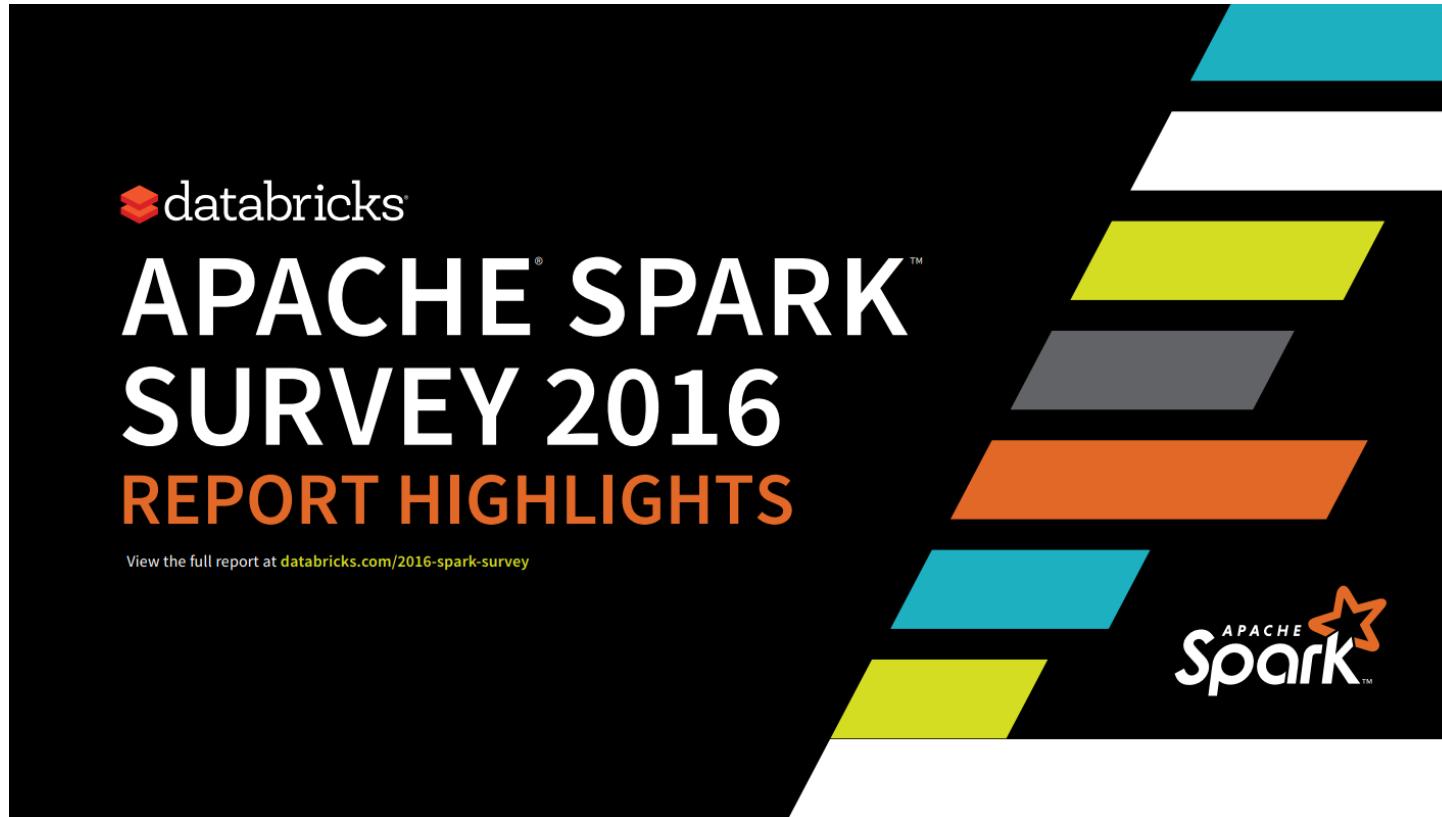
<http://ampcamp.berkeley.edu/wp-content/uploads/2012/06/dilip-joseph-amp-camp-2012-conviva-using-spark.pdf>

65 milioni di video stream / giorno

<https://www.qubole.com/blog/apache-spark-use-cases/>

Chi usa Spark e perché

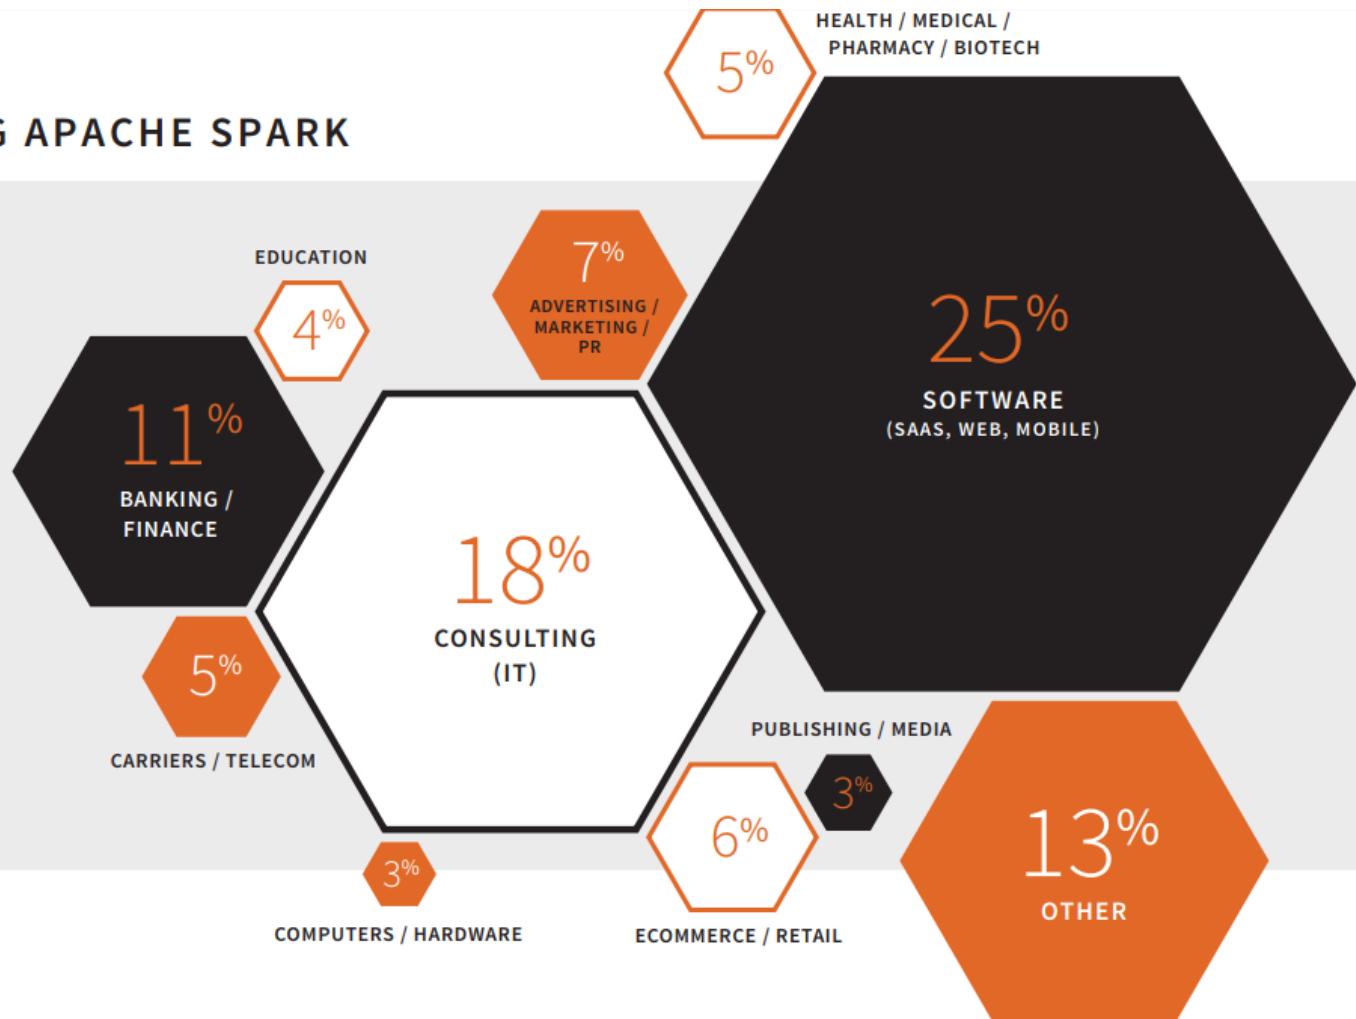
Nel 2016 i creatori di Spark hanno organizzato un sondaggio fra gli utenti.
(1615 risposte da parte di più di 900 organizzazioni)



http://pages.databricks.com/rs/094-YMS-629/images/2016_Spark_Infographic.pdf

Chi usa Spark e perché

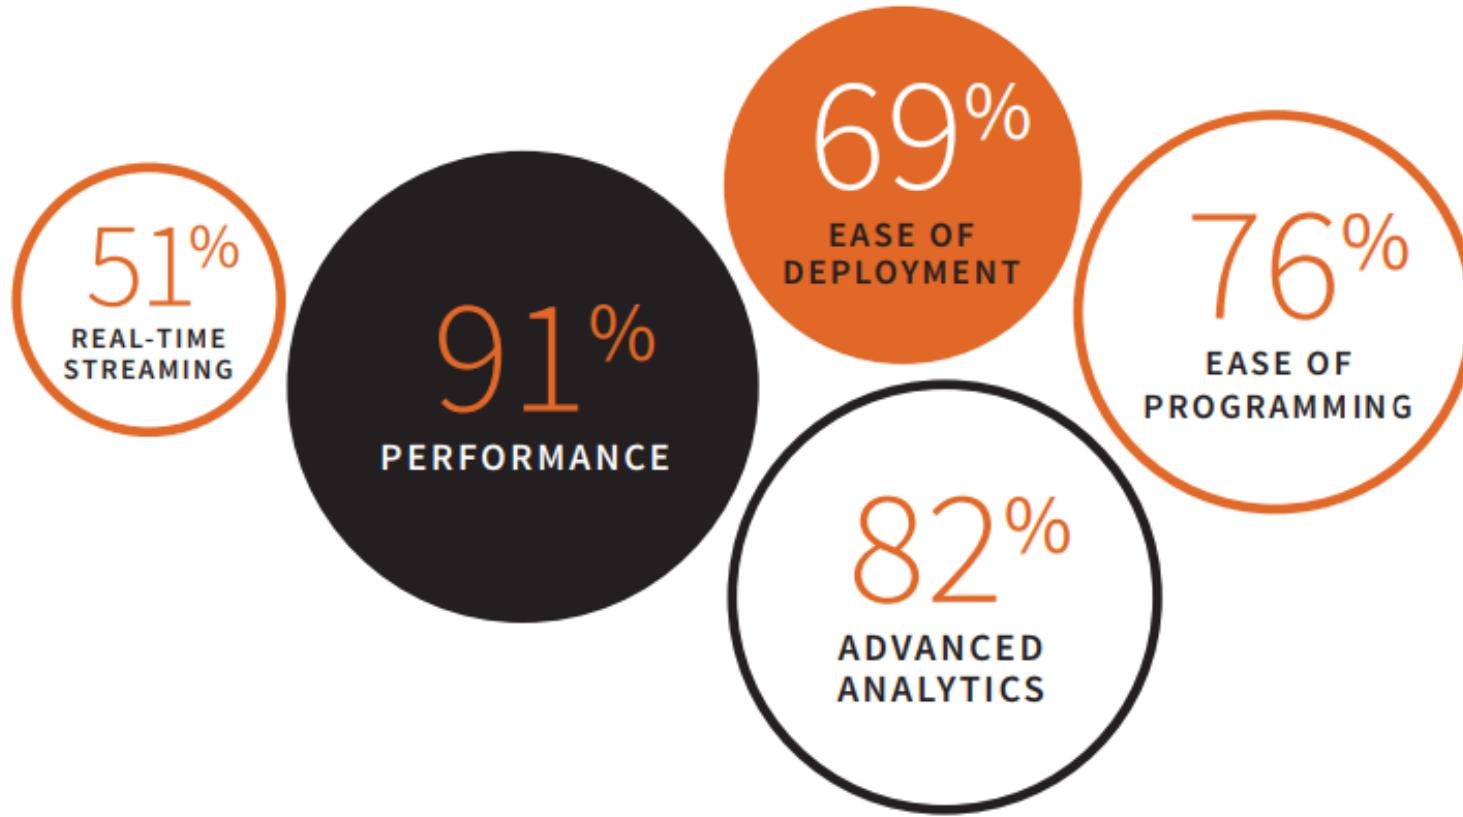
KEY INDUSTRIES USING APACHE SPARK



Chi usa Spark e perché

FEATURES USERS CONSIDER IMPORTANT

Respondents were allowed to select more than one feature.



Perché Spark?

Principali punti di forza

- Veloce
- Agnostico rispetto alla piattaforma
- Versatile
- Facile
- Capace di scalare
- Le «3 S» di Spark:
 - **Speed** = Velocità
 - **Simplicity** = Facilità di uso
 - **Support** = Vari linguaggi e sistemi di storage;
Provider commerciali + comunità grande, attiva e internazionale

<https://bigdata-madesimple.com/what-is-3ss-of-spark-and-its-effect-on-big-data/>

Punti di forza: Velocità

- Ingegnerizzato dal basso per ottenere le migliori performance
- 100x più veloce di *Hadoop MapReduce* grazie all'elaborazione in memoria (e altre ottimizzazioni)
- Veloce anche su disco.
- Record del mondo nel 2014 e 2016 per ordinamento su disco di grandi basi dati

New CloudSort Benchmark

Cost to sort 100TB of data



<https://databricks.com/blog/2016/11/14/setting-new-world-record-apache-spark.html>

Punti di forza: Facilità d'uso

- Approccio dichiarativo di alto livello ed API per agire su grandi datasets
- >100 operatori per trasformare data
- API basate su *DataFrame* e *Dataset* che permettono di manipolare facilmente dati strutturati e semi-strutturati

```
val errorCount = spark.read
    .text(fn)
    .filter($"value".like("%ERROR%"))
    .count()
```

Punti di forza: Facilità d'uso

WordCount in *MapReduce* (Java)

```
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;

public class WordCount {

    public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                word.set(tokenizer.nextToken());
                output.collect(word, one);
            }
        }
    }

    public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
        public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
            int sum = 0;
            while (values.hasNext()) {
                sum += values.next().get();
            }
            output.collect(key, new IntWritable(sum));
        }
    }

    public static void main(String[] args) throws Exception {
        JobConf conf = new JobConf(WordCount.class);
        conf.setJobName("WORDCOUNT");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        conf.setMapperClass(Map.class);
        conf.setCombinerClass(Reduce.class);
        conf.setReducerClass(Reduce.class);
        conf.setInputFormat(TextInputFormat.class);
        conf.setOutputFormat(TextOutputFormat.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
        JobClient.runJob(conf);
    }
}
```

WordCount in *Spark* (Scala)

```
val file = spark.textFile("hdfs://...")
val counts = file
    .flatMap(line => line.split(" "))
    .map(word => (word, 1)).reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
```

Punti di forza: Facilità d'uso

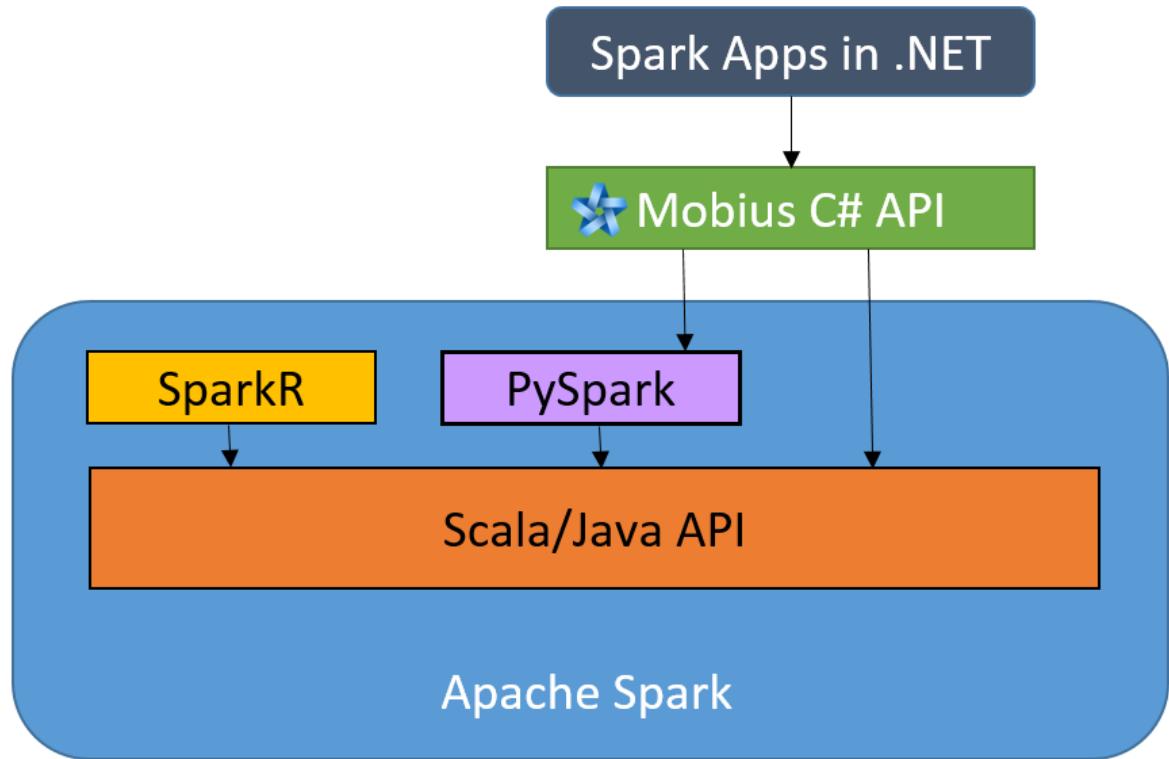
- Può essere eseguito interattivamente dalla shell o utilizzando un *NoteBook*
- Supporto per diversi linguaggi: **Scala**, **Java**, **Python**, **R**, **SQL**



E stanno arrivando altri linguaggi!

(progetti di terze parti)

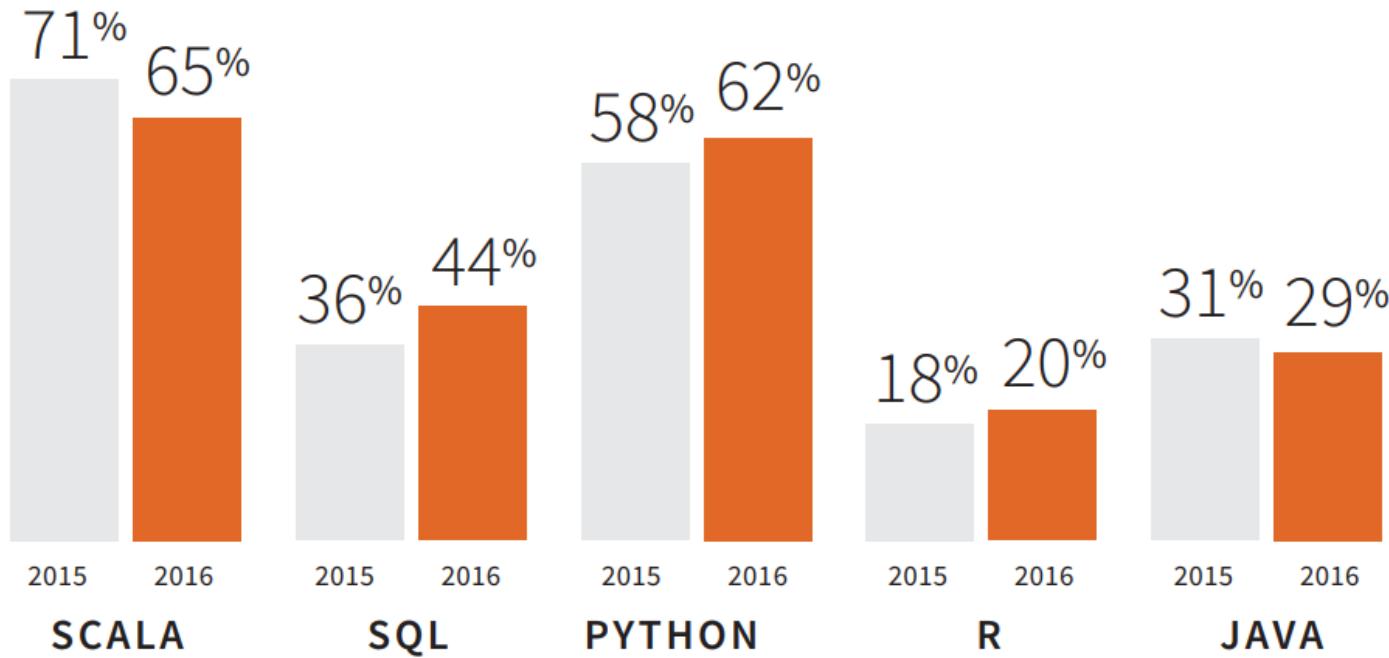
- *Closure*
- *C#/.net (progetto Möbius)*
- *Groovy*
- *Julia*



Quale linguaggio usare?

LANGUAGES USED IN APACHE SPARK

Respondents were allowed to select more than one language.



http://pages.databricks.com/rs/094-YMS-629/images/2016_Spark_Infographic.pdf

Quale linguaggio usare?

Java



PRO:

- Potente e completo
- Molto diffuso, conosciuto e utilizzato: probabilmente il linguaggio più usato per i Big Data
- Alte prestazioni

CONTRO:

- Verboso
- Non supporta *Read-Evaluate-Print-Loop*

Quale linguaggio usare?

R



PRO:

- Il linguaggio di riferimento per Statistici e Analisti di Dati
- Buone librerie di visualizzazione

CONTRO:

- Non è considerato un linguaggio di utilizzo generale
- Prestazioni non alte
- C'è un certo ritardo nell'integrazione con Spark

Quale linguaggio usare?

SQL



PRO:

- Utilizzabile anche da non programmati
- Trend di utilizzo in crescita

CONTRO:

- Inadatto per alcuni tipi di applicazioni

Quale linguaggio usare?

Python



PRO:

- Linguaggio molto noto e diffuso, anche nel mondo dei Big Data
- Facile da imparare
- Molto conciso ed espressivo
- Buone librerie di visualizzazione

CONTRO:

- Più lento
- «*Duck typing*». Controllo sui tipi a run-time

Quale linguaggio usare?

Scala



PRO:

- Spark è scritto in Scala: ogni feature del framework è immediatamente disponibile
- Compatibile con le librerie Java
- Controllo sui tipi durante la compilazione
- Alte prestazioni

CONTRO:

- Meno diffuso di Java o Python
- Marginalmente meno facile di Python

In questo corso useremo:



```
object HelloWorld {  
    def main(args: Array[String]): Unit = {  
        println("Hello, world!")  
    }  
}
```

Spark ha qualche punto debole

- Non ha un suo sistema di storage
- Grande consumo di risorse (in particolare la memoria) \Rightarrow Spark è dispendioso
- Può gestire anche dataset molto grandi, ma richiede un po' di attenzione.
- Non ha supporto per il *real-time processing* (*Spark Streaming* usa micro-batches)
- Tende a generare tanti file relativamente piccoli. Questo è un male, ad es. per *Hadoop HDFS* o *Amazon S3*.
- Può richiedere ottimizzazione manuale
- Ha una latenza considerevole

Non solo Apache Spark!

Alcune alternative:



(creato da Twitter)



Apache Flink



IBM InfoSphere Streams

Il contesto (e un po' di storia)

Contesto

La potenza di calcolo ha continuato a crescere ininterrottamente per ottant'anni



Univac I (1951)
~ 1 000 000\$
1905 operazioni al secondo



Cray-1 (1976)
~ 10 000 000\$, 115 Kw
1.9 GFLOPS



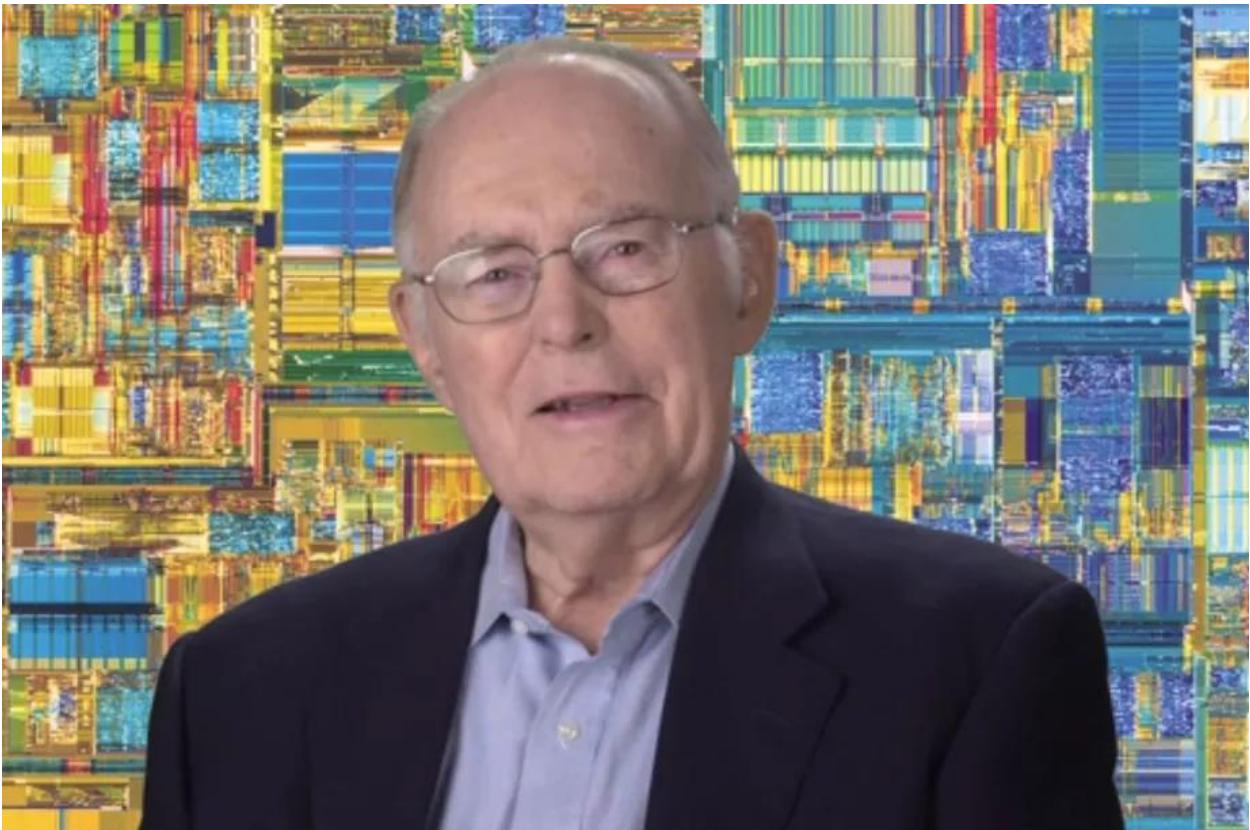
iPad 2 (2011)
<1 000\$
1.6 GFLOPS

La legge di Moore



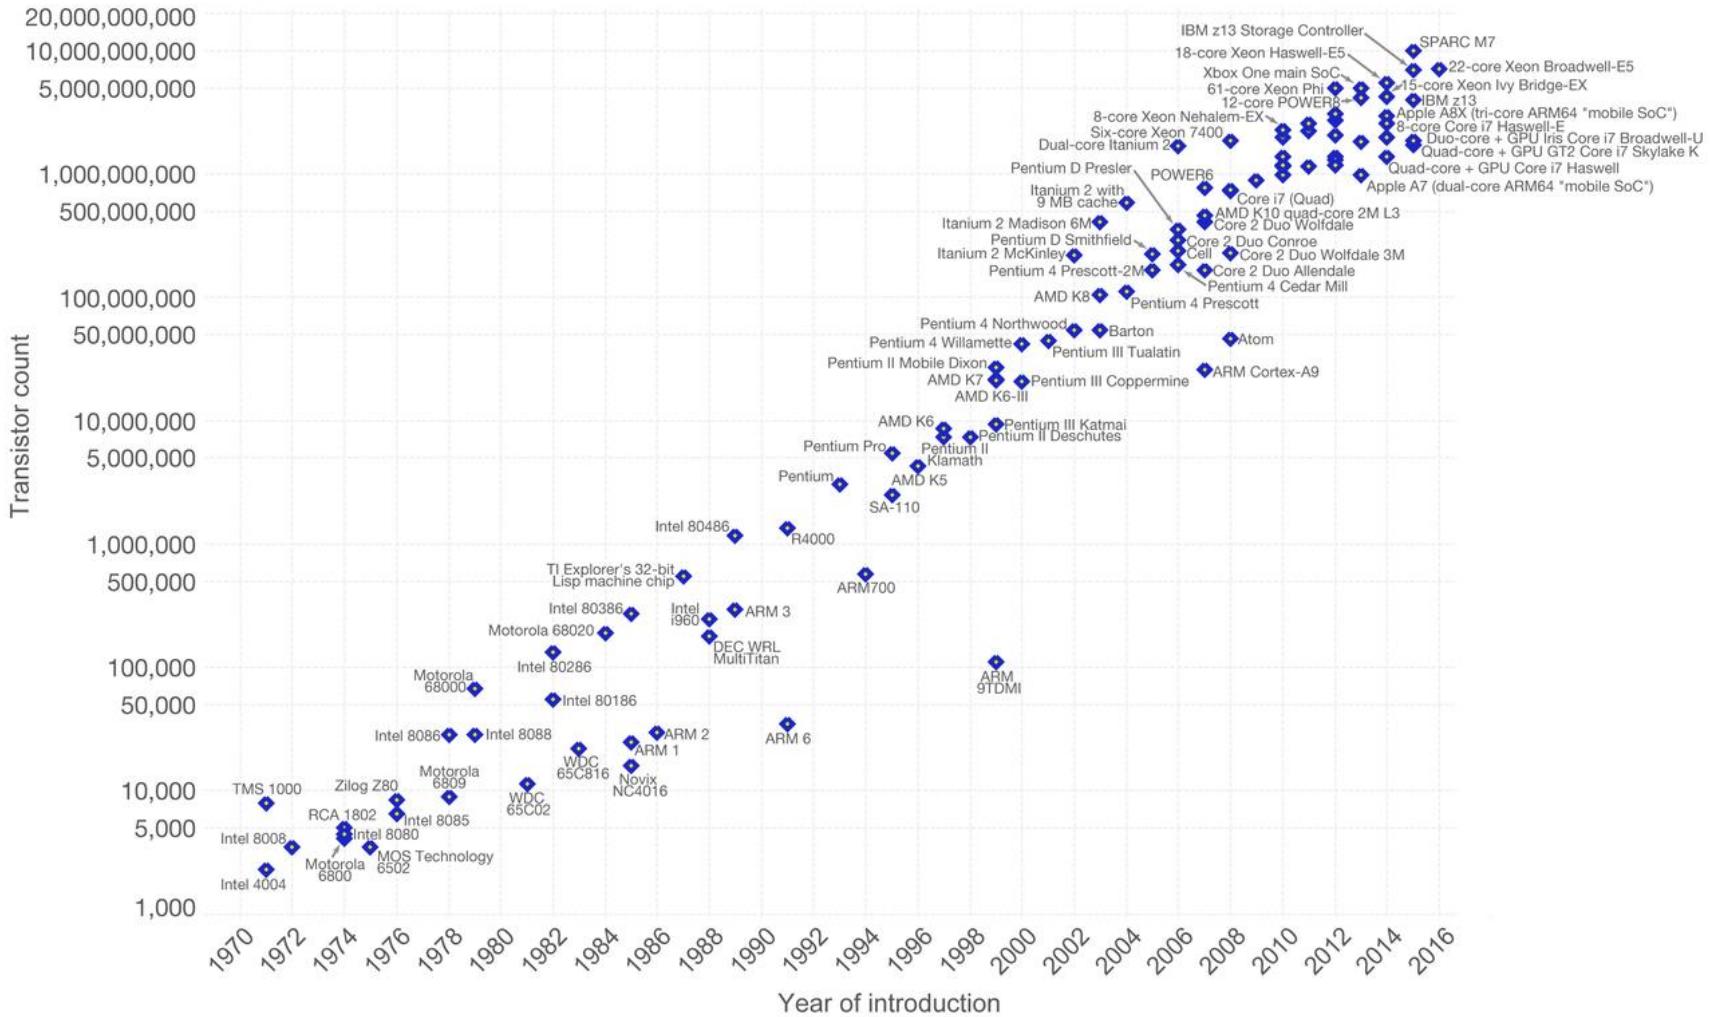
Gordon Earle Moore, nato nel 1929

Co-fondatore di Intel



Moore's Law – The number of transistors on integrated circuit chips (1971-2016)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

The data visualization is available at OurWorldinData.org. There you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

La legge di Moore è ancora valida?

DEATH NOTICE

Long beloved by both engineers and computer scientists because of ongoing performance benefits ceaselessly and seemingly effortlessly achieved. From the age of fifty, Moore's Law began to age rapidly, passing into senescence and then, at the beginning of this month, into oblivion. Moore's Law leaves a thriving multi-trillion dollar global semiconductor and electronics industry, along with a growing set of questions about how that industry will survive its passing. In lieu of flowers, donations should be lavished on Intel shares.

https://www.theregister.co.uk/2018/01/24/death_notice_for_moore_s_law/

Un'altra strategia per andare veloci: il parallelismo

Largest
Amazon EC2
instance: 128
virtual CPUs

Azul Systems Vega 3
Cores per chip: 54
Cores per system: 864

The Future is Parallel

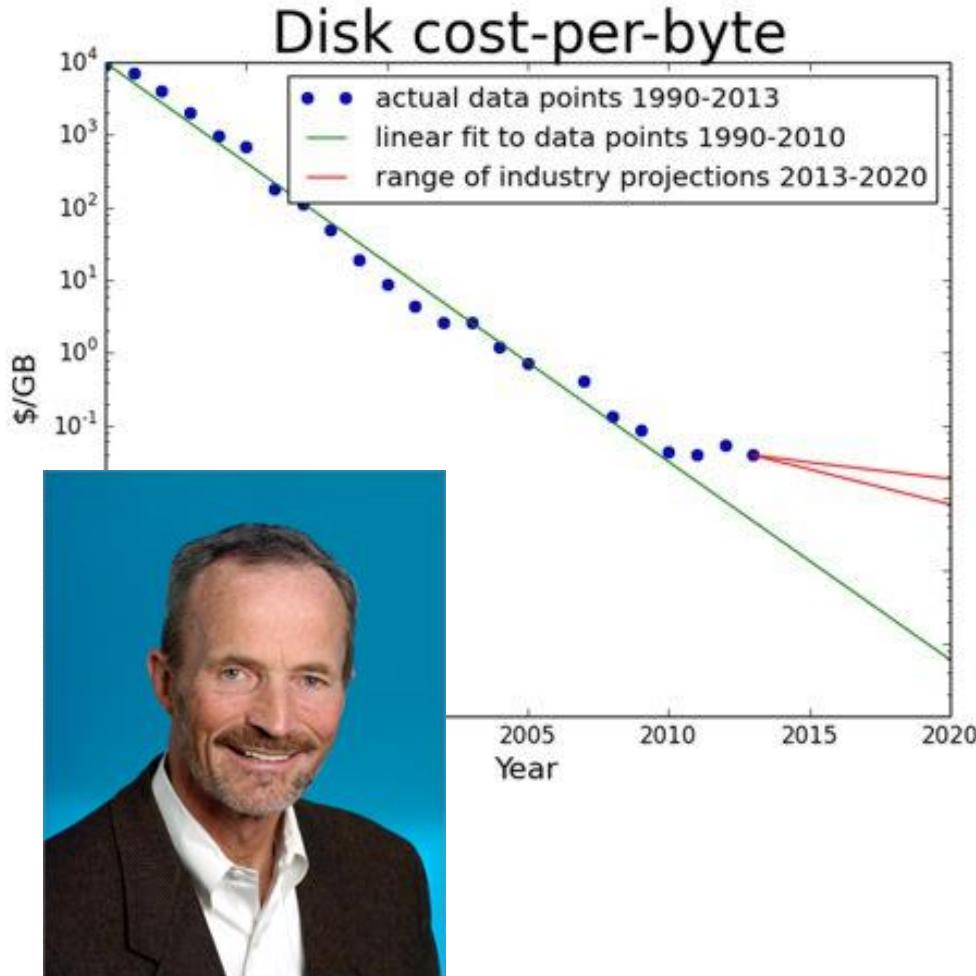
Intel Xeon
24 cores
48 threads

AMD
Opteron
16 cores

Tilera Gx-
3000
100 cores

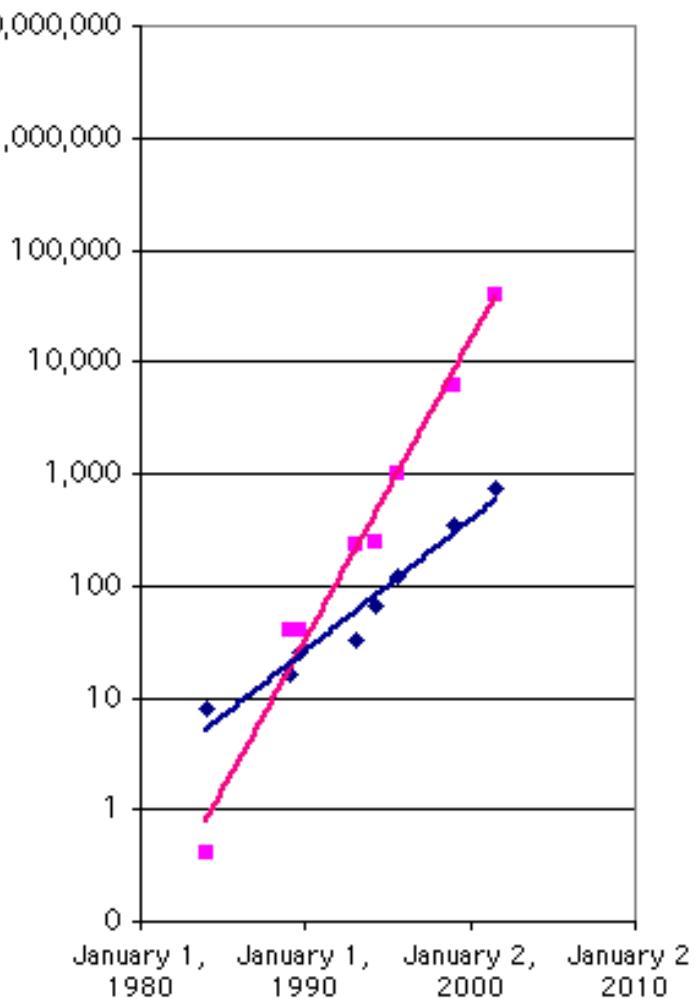
La legge di Kryder

Le prestazioni dei dischi crescono più velocemente della potenza di calcolo



Mark H. Kryder

Legge di Kryder (viola)
vs.
Legge di Moore (blu)



Le prestazioni dei dischi crescono più velocemente della potenza di calcolo

Hardware Trends

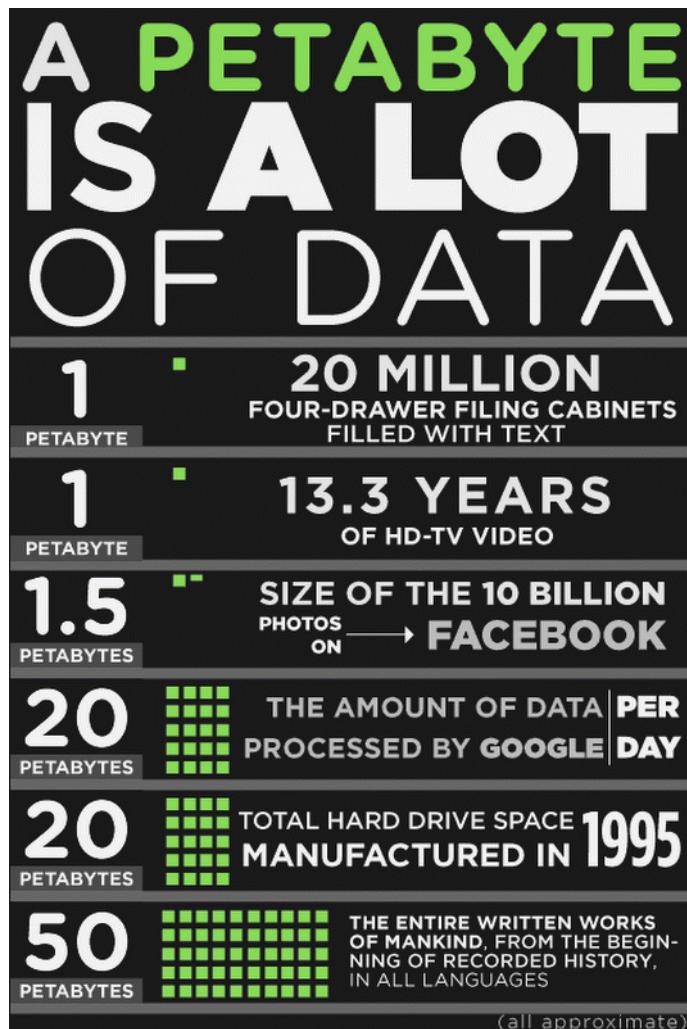
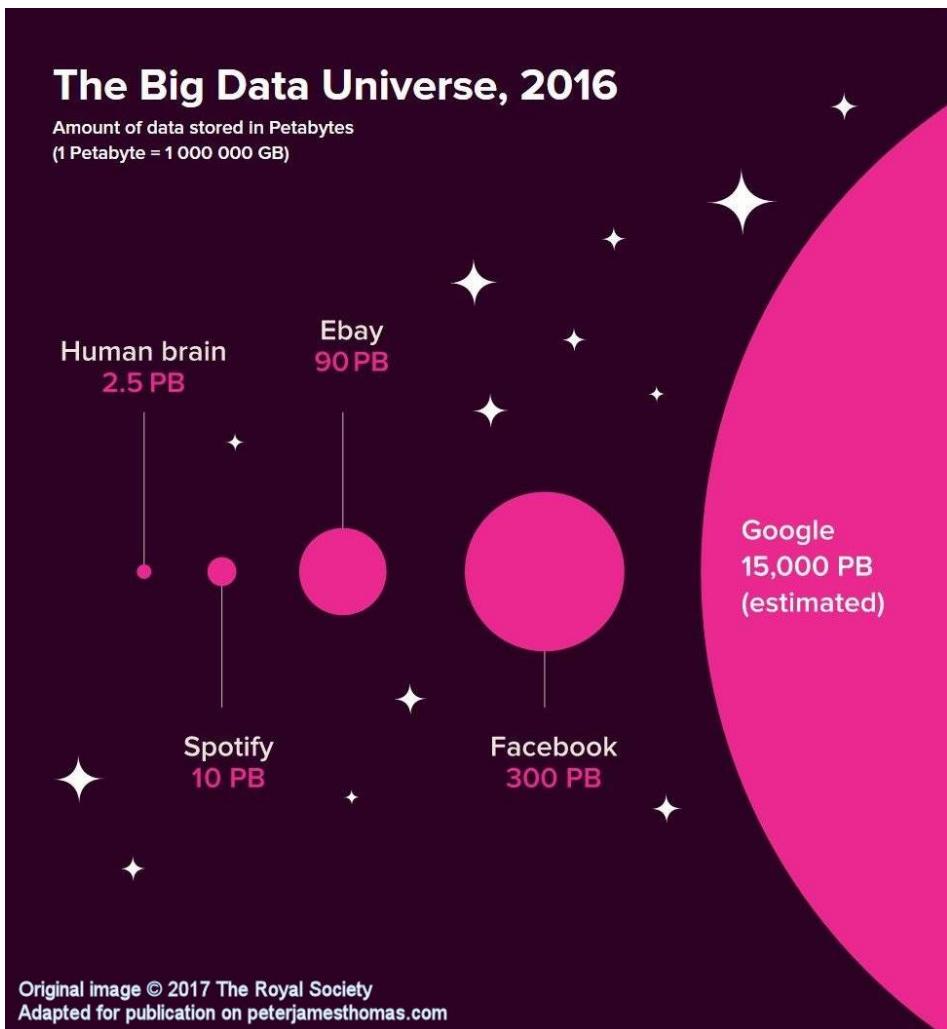
| | 2010 | 2015 | |
|---------|------------------|-------------------|-----|
| Storage | 50+MB/s (HDD) | 500+MB/s (SSD) | 10X |
| Network | 1Gbps | 10Gbps | 10X |
| CPU | ~3GHz | ~3GHz | :(|

L'avvento dei BIG Data

BIG DATA:

- 500 milioni di tweets al giorno
- Traffico di quasi 2 GB/mese sui telefoni cellulari
- Amazon vende 600 prodotti al secondo
- 200 miliardi di email vengono scambiate ogni giorno
- MasterCard esegue 74 miliardi di transazioni ogni anno
- Ci sono quasi 6000 voli di linea al giorno.

Quanto sono «BIG» i Big Data?

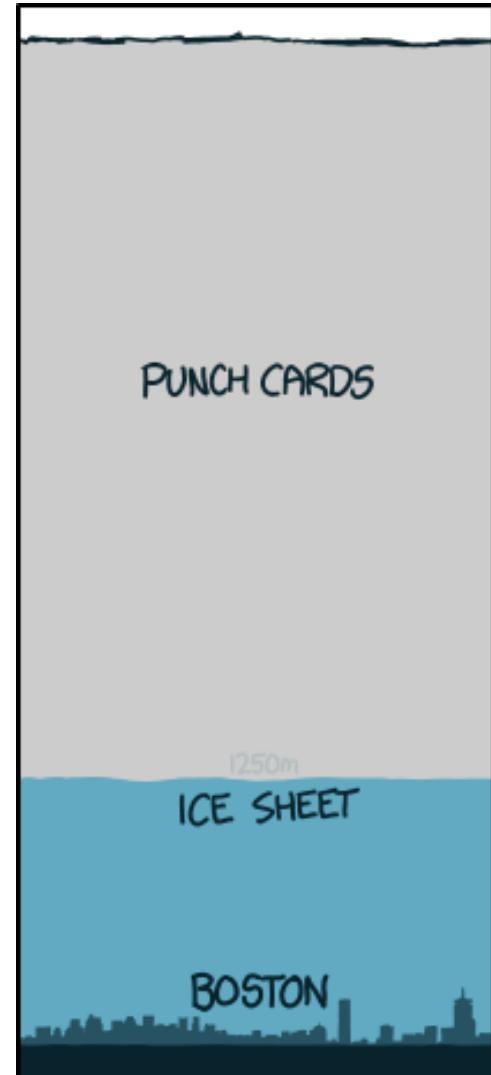


Intermezzo : aspetta... quanti dati ha Google???

“If all digital data were stored on punch cards,
how big would Google’s data warehouse be?”



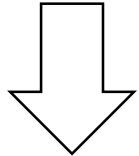
Randall Munroe



<https://blog.ted.com/using-serious-math-to-answer-weird-questions-randall-munroe-at-ted2014/>

Riassumendo

- La velocità di calcolo cresce, ma non tanto quanto vorremmo
- Lo storage cresce molto più velocemente
- Con uno storage economico e virtualmente illimitato la quantità di dati raccolti diventa terrificante



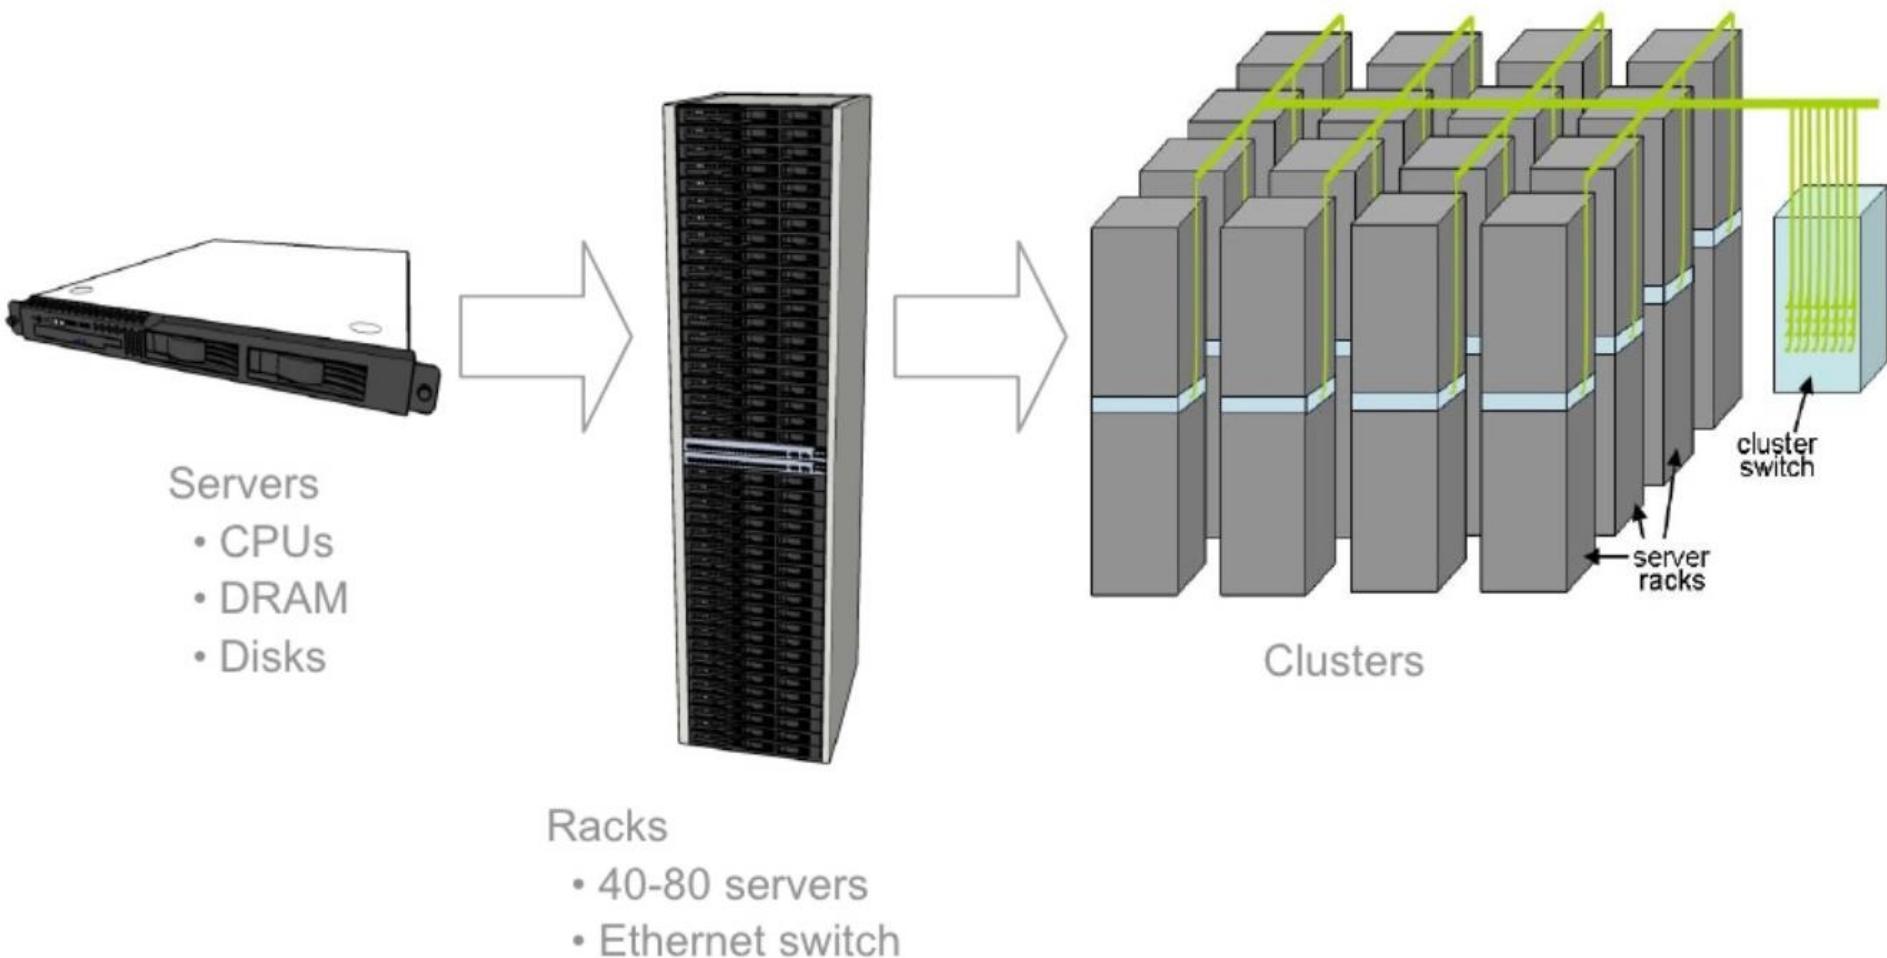
Il modo più efficace di analizzare questa enorme massa di dati è il *cluster computing*

Esistono cluster con centinaia e anche migliaia di nodi!

Come si programma questa roba?



Come si programma questa roba?



Nel primo decennio del secolo Google pubblica tre articoli fondamentali

2003

GFS - Google File System

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung
Google

ABSTRACT

We have designed and implemented the Google File System (GFS) to meet the rapidly growing demands of Google's distributed data processing system. GFS is a fault-tolerant file system running on inexpensive commodity hardware, and it delivers high bandwidth access to large amounts of data.

While sharing many of the design goals as previous distributed file systems, our design has been driven by characteristics of Google's data processing needs. In particular, both current and anticipated, that reflect a scaled deployment of thousands of machines, and a large number of clients.

In this paper we present our file system interface extensions designed to support distributed applications. We discuss many aspects of the system, including lessons learned from both micro-benchmarks and real world use.

Categories and Subject Descriptors

D.4.2 - Distributed file systems

General Terms

Design, scalability, performance, measurement

Keywords

Fault tolerance, scalability, data storage, clustered storage

The authors can be reached at the following addresses:

[sanjay, gobi, shuleung]@google.com

Premise to make digital a last copy of all of part of the work is served or not. This is to prevent the provided that copies are not made or distributed for profit or commercial advantage and the copy is not changed in any way. The digital version of the article is provided for personal research only and is not intended for sale. It may be copied and given to your colleagues; however, to re-use more than the permitted number of copies or to republish, to post on servers or to redistribute to lists, requires prior specific permission and payment to the copyright holders.

Fourth, co-designing the applications and the file system

API benefits the overall system by increasing our flexibility.

Copyright 2003 ACM 1511-3350/03/0705-0001 \$15.00

2004

MapReduce

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat
jeff@google.com, sanja@google.com
Google, Inc.

Abstract

MapReduce is a programming model and an associated infrastructure for processing large amounts of data held in a distributed system of thousands of machines. The model has been driven by key observations of our application workloads and their required semantics, both current and anticipated. One is that often parallel algorithms are better suited to distributed file systems than to shared memory. This has led us to re-examine traditional choices and explore radically different ones.

The file system has successfully met our storage needs. It is reliable, fast, and provides a simple interface for the generation and processing of data used by our service as well as research and development efforts that require fast access to large amounts of data.

Our system has been used to process tens of thousands of terabytes of storage across thousands of disks on over 1000 machines, and to support a wide range of applications run by hundreds of clients.

In this paper we present our system's interface extensions designed to support distributed applications. We discuss many aspects of the system, including lessons learned from both micro-benchmarks and real world use.

As a reaction to this complexity, we designed a new abstraction that allows us to express the simple computations we were trying to perform but hides the messy details of partitioning the input data, scheduling the program's execution across a set of machines, handling network communication, and performing fault recovery.

This allows programmers to focus on the logic of their computation and ignore the details of partitioning the input data, scheduling the program's execution across a set of machines, handling network communication, and performing fault recovery.

Our implementation of MapReduce runs on a large cluster of commodity machines and is highly scalable. It typically highlights the strengths of parallel processing systems of data on thousands of machines. Programmers can take advantage of the built-in fault tolerance guarantees provided by MapReduce and the fact that one of these standard MapReduce jobs are executed on Google's clusters every day.

1 Introduction

Over the past few years, the authors and many others at Google have implemented hundreds of special-purpose computations that process large amounts of raw data. These computations are typically used in search engines, news crawlers, and other distributed systems. Some may be implemented on a single machine, while others are distributed on another, whether simultaneously or later in time. Given the variety of these applications, it is important to have a set of performance optimization and latency guarantees, as well as a simple interface for distributed applications.

Fourth, co-designing the applications and the file system

API benefits the overall system by increasing our flexibility.

To appear in OSDI 2004

2006

BigTable

Bigtable: A Distributed Storage System for Structured Data

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Heis, Deborah A. Wallach
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber
[fay, jeff, sanja, wilson, deborah, mike, tushar, andrew, rgruber]@google.com

Google, Inc.

Abstract

Bigtable is a distributed storage system for managing structured data. It is designed for data sets that are too large and the computations have to be distributed across hundreds or thousands of machines in order to finish in reasonable time. Bigtable provides a simple interface to enable the computation, distribute the data, and handle failures. It also provides a mechanism for clients to observe the original simple computation and to react to changes in the data set with which these issues.

As a reaction to this complexity, we designed a new abstraction that allows us to express the simple computations we were trying to perform but hides the messy details of partitioning the input data, scheduling the program's execution across a set of machines, handling network communication, and performing fault recovery.

This allows programmers to focus on the logic of their computation and ignore the details of partitioning the input data, scheduling the program's execution across a set of machines, handling network communication, and performing fault recovery.

Our implementation of MapReduce runs on a large cluster of commodity machines and is highly scalable. It typically highlights the strengths of parallel processing systems of data on thousands of machines. Programmers can take advantage of the built-in fault tolerance guarantees provided by MapReduce and the fact that one of these standard MapReduce jobs are executed on Google's clusters every day.

The major contributions of this work are a simple and powerful programming model for distributed computation and distribution of large-scale computations, combined with an implementation of this model that runs on commodity hardware (mostly PCs).

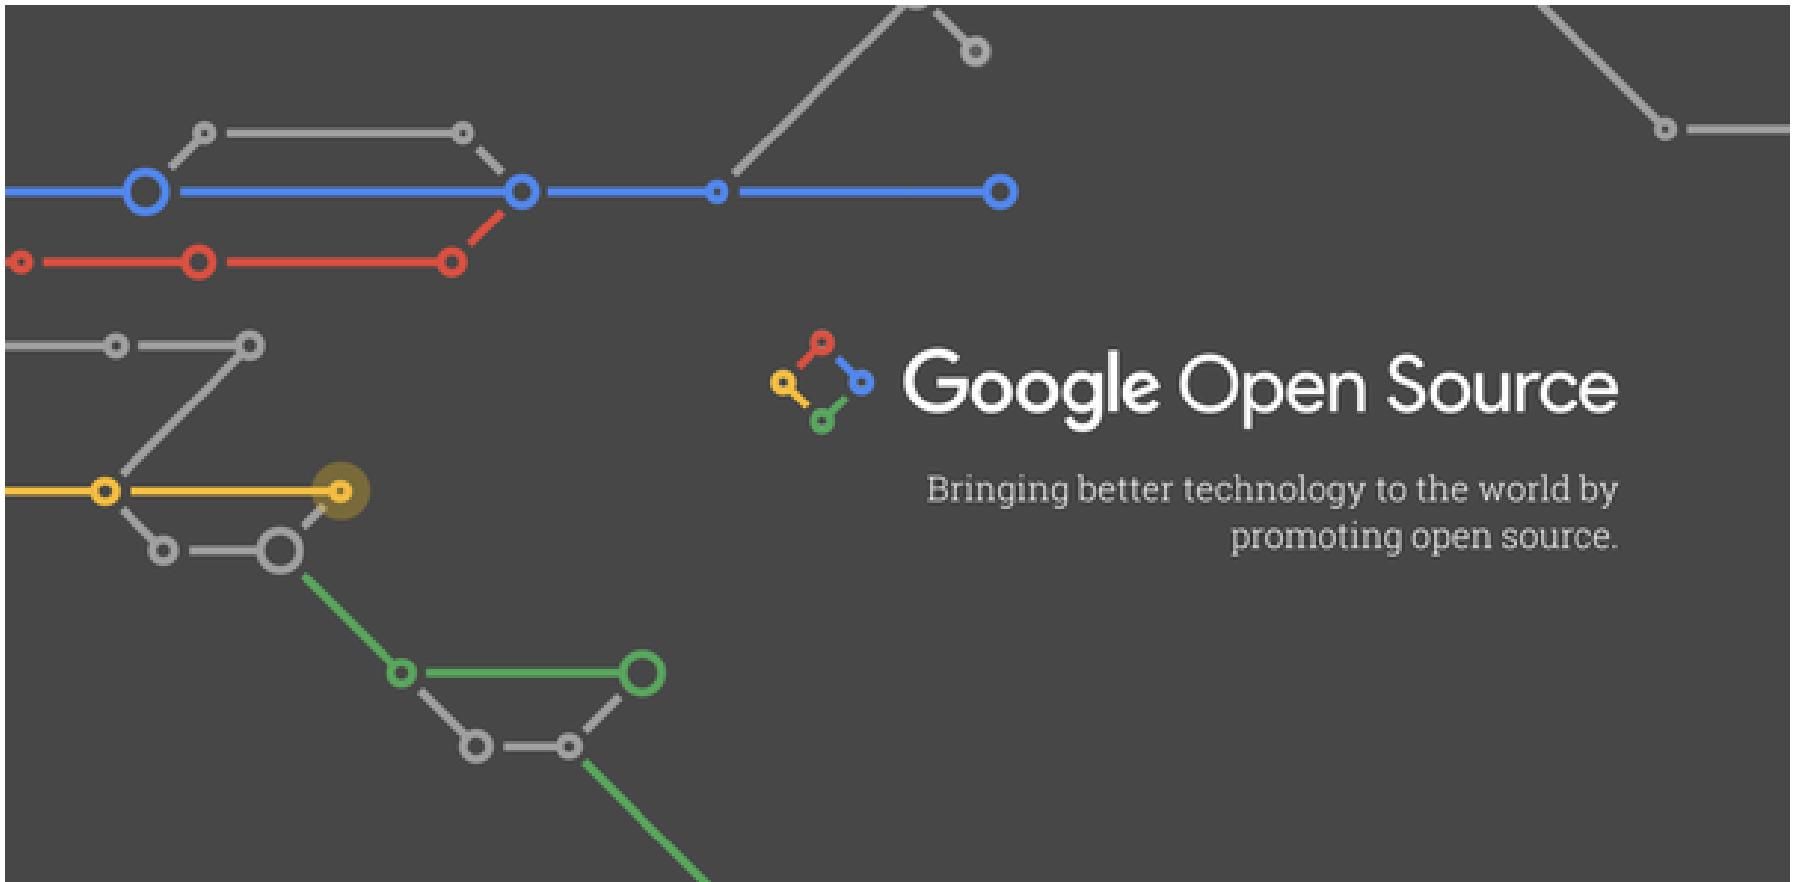
Section 2 describes the basic programming model and gives several examples. Section 3 describes the implementation of the MapReduce model and its use in our cluster-based computing environment. Section 4 describes the system as it is used, how it has been used, and what we have found useful. Section 5 has performance measurements of our implementation for a variety of workloads. Section 6 discusses the use of MapReduce in Google including our experiences in using it as the basis

To appear in OSDI 2006

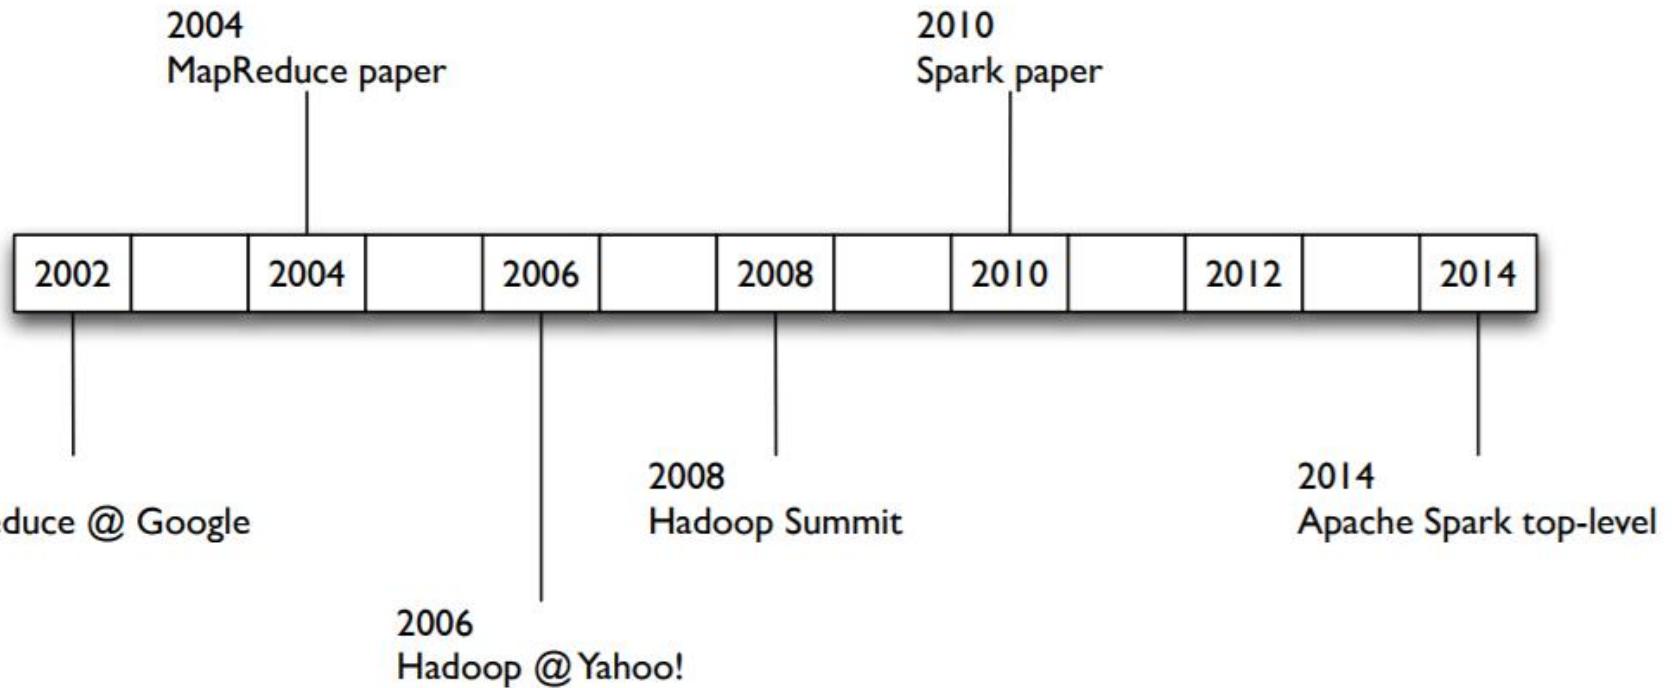
1

ma non pubblica il codice dell'implementazione!

... poi (un po' in ritardo) abbraccia l'Open Source



Nel frattempo nascono Hadoop e Spark



Subito prima della nascita di Spark

Hadoop è la soluzione di riferimento per i Big Data:

- HDFS per lo storage e
- MapReduce per il calcolo

MapReduce ha delle limitazioni: è relativamente lento e difficile da programmare.

Inoltre gli utenti cominciano a richiedere:

- Algoritmi più complessi, basati su passi multipli
- Query ad-hoc interattive
- Processing in real time

Queste feature sono difficilmente implementabili in MapReduce

Subito prima della nascita di Spark

Per ovviare a queste limitazioni vengono sviluppati dei sistemi *specializzati*, uno per ogni tipo di esigenza:

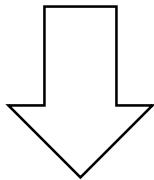
- Pregel
- Giraph
- Dremel
- Drill
- Impala
- Presto
- Storm
- S4
-



Subito prima della nascita di Spark

Problemi con i sistemi specializzati

- Molti sistemi da controllare, configurare e installare
- Difficile combinarli fra loro anche se spesso servirebbe (ad es. leggere dati con SQL e poi usare degli algoritmi di machine learning)
⇒ In molti casi il trasferimento dei dati fra due sistemi diventa un costo dominante



Spark nasce come un engine unificato

2009 : Nascita di Spark

Un gruppo di ricerca a UC Berkeley

Il gruppo lavora su un framework per il controllo di cluster che possa sostenere diversi tipi di cluster computing system. Spark nasce come esempio.



An Architecture for Fast and General Data Processing
on Large Clusters

All'inizio focalizzato sul *machine learning*, poi diventa uno strumento generale



Matei Zaharia

Electrical Engineering and Computer Sciences
University of California at Berkeley

Technical Report No. UCB/EECS-2014-12
<http://www.eecs.berkeley.edu/Pubs/TechRpts/2014/EECS-2014-12.html>

February 3, 2014

<https://www.kdnuggets.com/2015/05/interview-matei-zaharia-creator-apache-spark.html>

2009 : Nascita di Spark

Un gruppo di ricerca a UC Berkeley

Due idee chiave:

- Tenere i risultati intermedi in memoria
- Fare uno strumento generale che possa essere utilizzato in contesti differenti (es. streaming e graph processing)

Nel 2013 si muove in Apache

Nel 2014 diventa un ***Top-Level Project*** (TLP)



An Architecture for Fast and General Data Processing
on Large Clusters

Matei Zaharia



Electrical Engineering and Computer Sciences
University of California at Berkeley

Technical Report No. UCB/EECS-2014-12
<http://www.eecs.berkeley.edu/Pubs/TechRpts/2014/EECS-2014-12.html>

February 3, 2014

<https://www.kdnuggets.com/2015/05/interview-matei-zaharia-creator-apache-spark.html>

2014: Spark è Apache Top Level Project



Diventare TPL è un processo complicato e laborioso.

Quando un progetto diventa TLP può contare su una comunità grande e attiva e su una grande visibilità

| How to plan your graduation | | | | | | |
|-----------------------------------|------------------------------------|-----|---------------|-----------------------------------------------------------------------------------------------------------------------------------|-----|-----|
| Sun | Mon | Tue | Wed | Thu | Fri | Sat |
| | | | | | | |
| community graduation vote | let vote run for at least 72 hours | | tally | | | |
| prepare charter in your community | | | | | | |
| IPMC graduation vote | let vote run for 72 hours | | tally | Use this time as a backup if/when the IPMC has questions or issues with the proposal. These can delay acceptance of your proposal | | |
| resolution to board@ | | | | | | |
| agenda fixed | | | board meeting | | | |
| | | | | | | |

A red tag is attached to the 'board meeting' cell in the fourth row, with the text '3rd wednesday of each month' written on it.

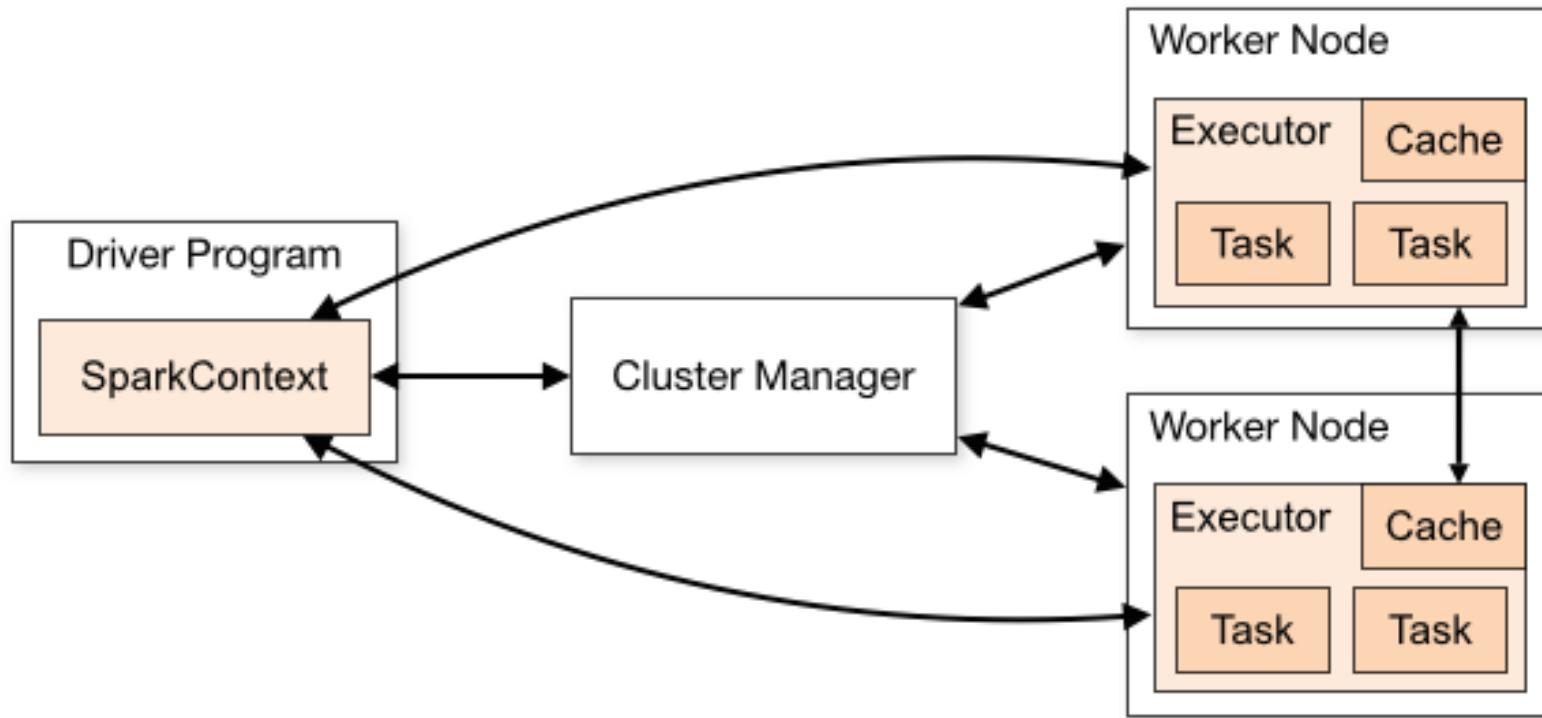
Apache is a well recognized brand



Spark,
com'è fatto?

Struttura di un programma Spark

L'elaborazione è distribuita su più nodi



L'ecosistema di Spark

Cluster manager



Storage



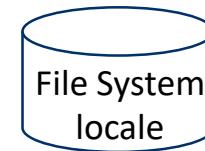
amazon
S3



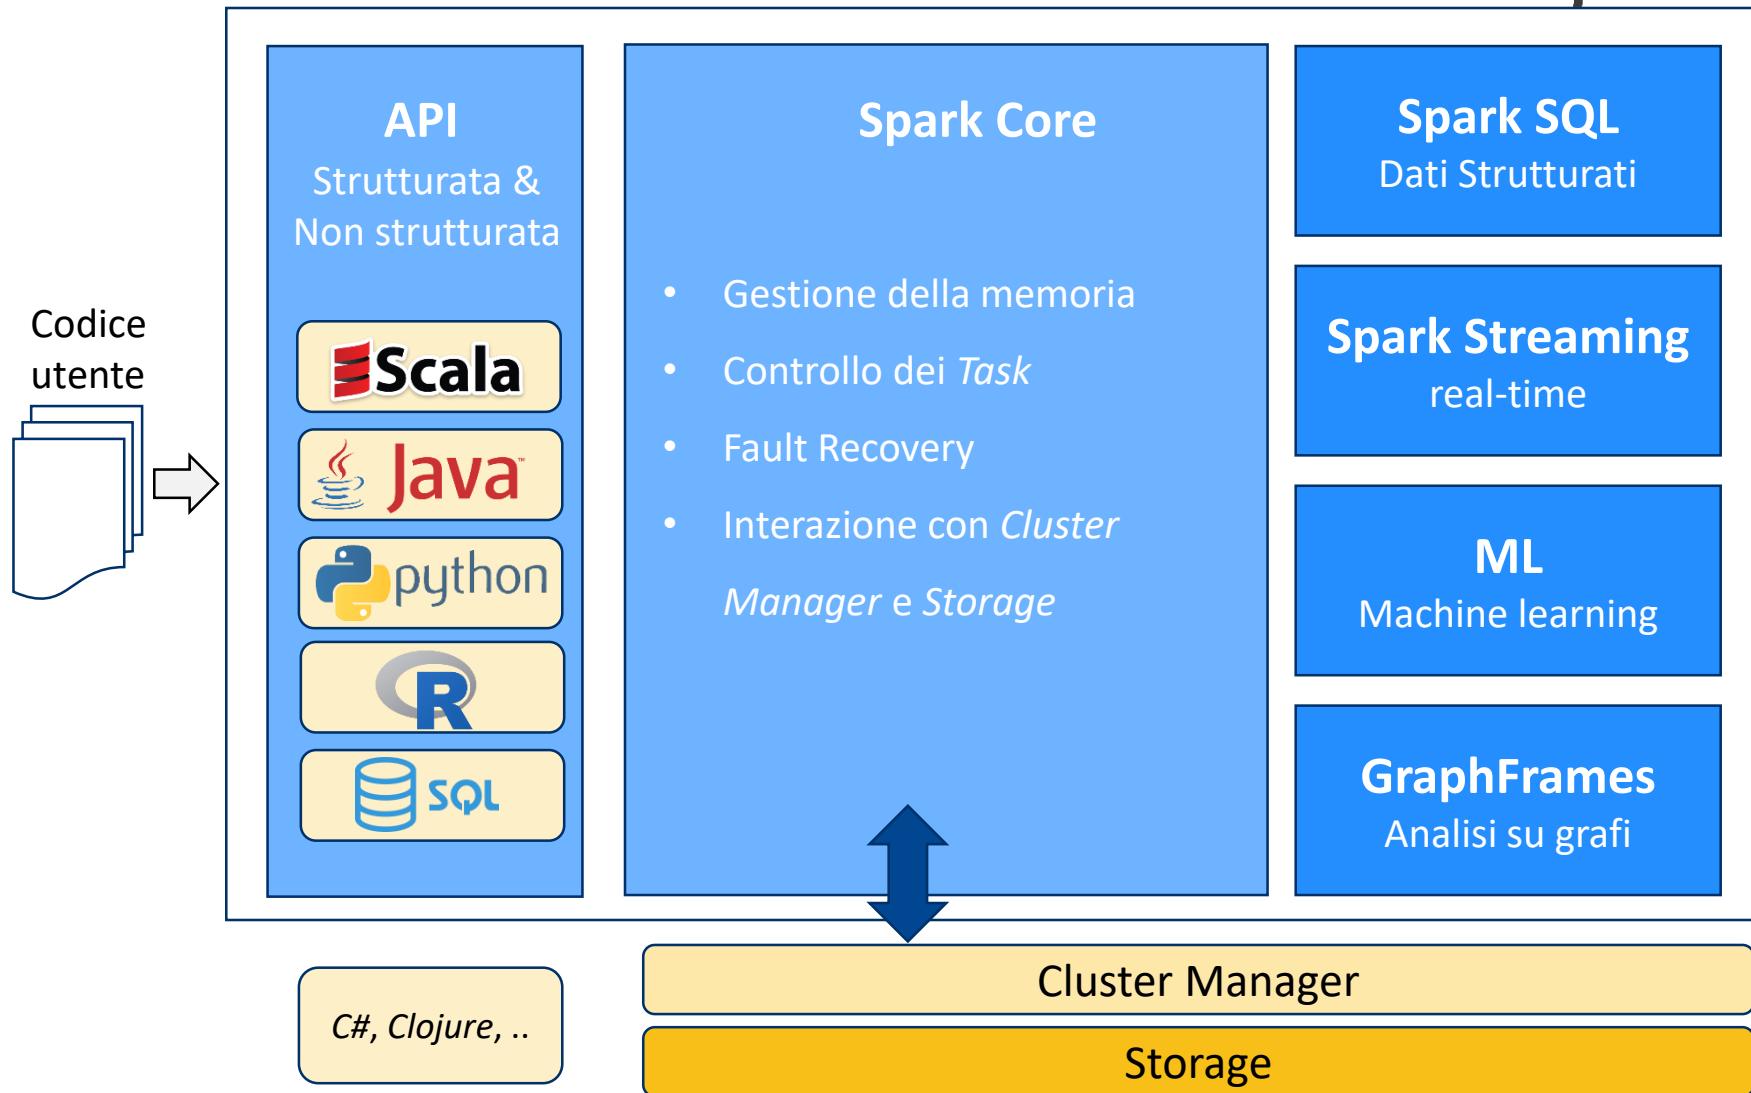
cassandra



Google Cloud Storage



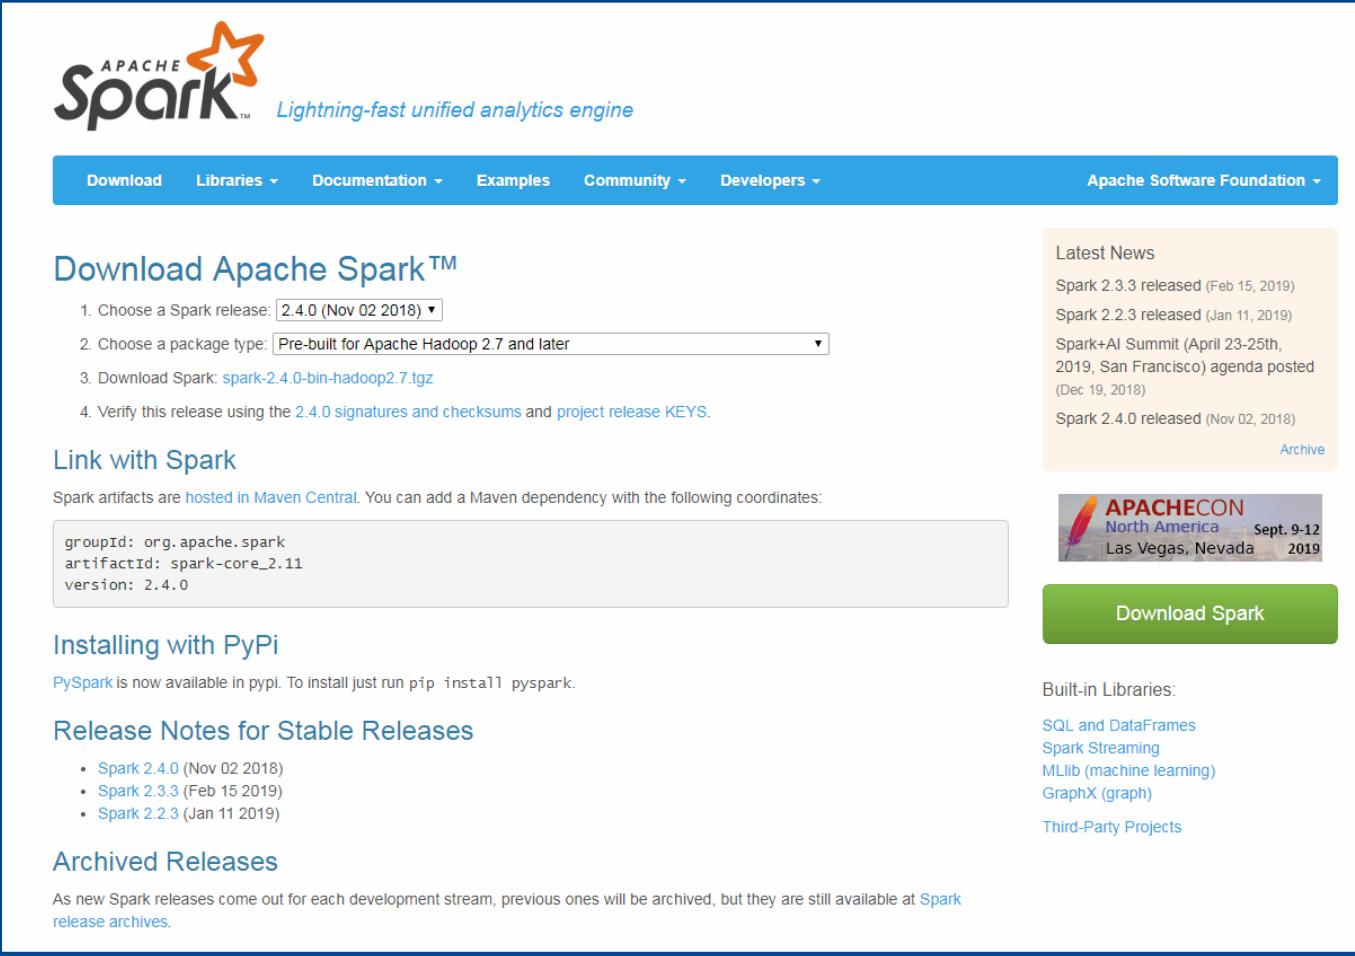
L'ecosistema di Spark



In pratica?

Per fare degli esperimenti a casa

<https://spark.apache.org/downloads.html>



The screenshot shows the Apache Spark downloads page. At the top, there's the Apache Spark logo with the tagline "Lightning-fast unified analytics engine". Below the logo is a navigation bar with links for Download, Libraries, Documentation, Examples, Community, Developers, and Apache Software Foundation. The main content area has a heading "Download Apache Spark™" and a numbered list of steps:

1. Choose a Spark release: 2.4.0 (Nov 02 2018) ▾
2. Choose a package type: Pre-built for Apache Hadoop 2.7 and later
3. Download Spark: spark-2.4.0-bin-hadoop2.7.tgz
4. Verify this release using the 2.4.0 signatures and checksums and project release KEYS.

Below this, there's a section titled "Link with Spark" with Maven dependency coordinates:

```
groupId: org.apache.spark  
artifactId: spark-core_2.11  
version: 2.4.0
```

There's also a section for "Installing with PyPi" and "Release Notes for Stable Releases" which lists releases from Nov 02 2018 to Jan 11 2019. A "Archived Releases" section notes that previous versions are available at the [Spark release archives](#). To the right, there's a "Latest News" sidebar with links to news items like "Spark 2.3.3 released (Feb 15, 2019)" and "Spark 2.2.3 released (Jan 11, 2019)". There's also a "APACHECON North America Sept. 9-12 Las Vegas, Nevada 2019" announcement and a large green "Download Spark" button. A sidebar on the right lists "Built-in Libraries" (SQL and DataFrames, Spark Streaming, MLlib, GraphX) and "Third-Party Projects".

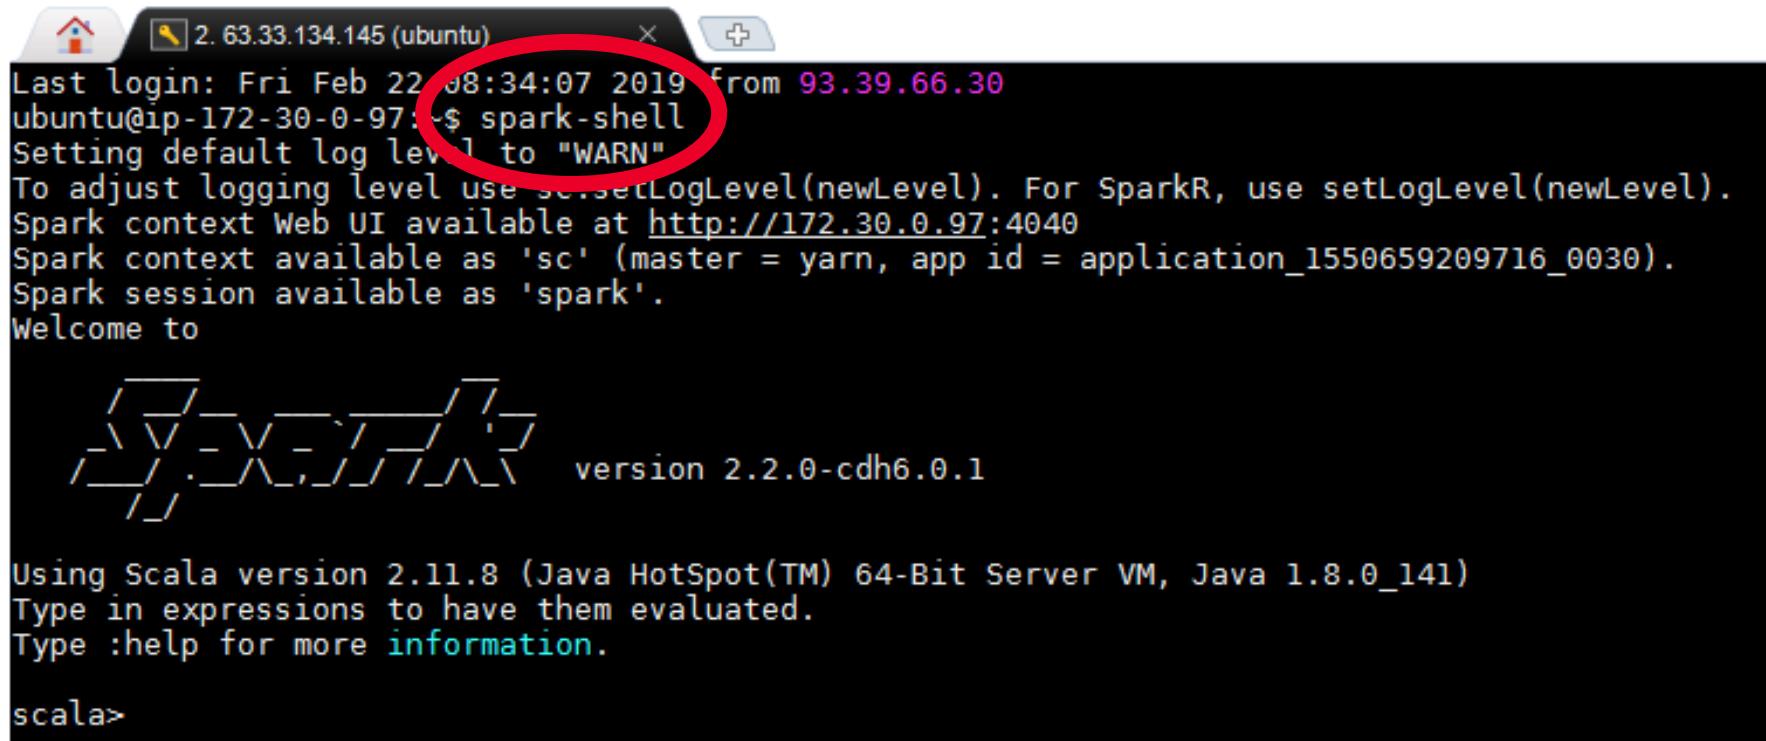
Per fare degli esperimenti a casa

Spark gira anche su un singolo computer

Lo stesso codice funziona su un cluster di
1000 nodi!

Come useremo Spark nel laboratorio di oggi?

La Shell di Spark



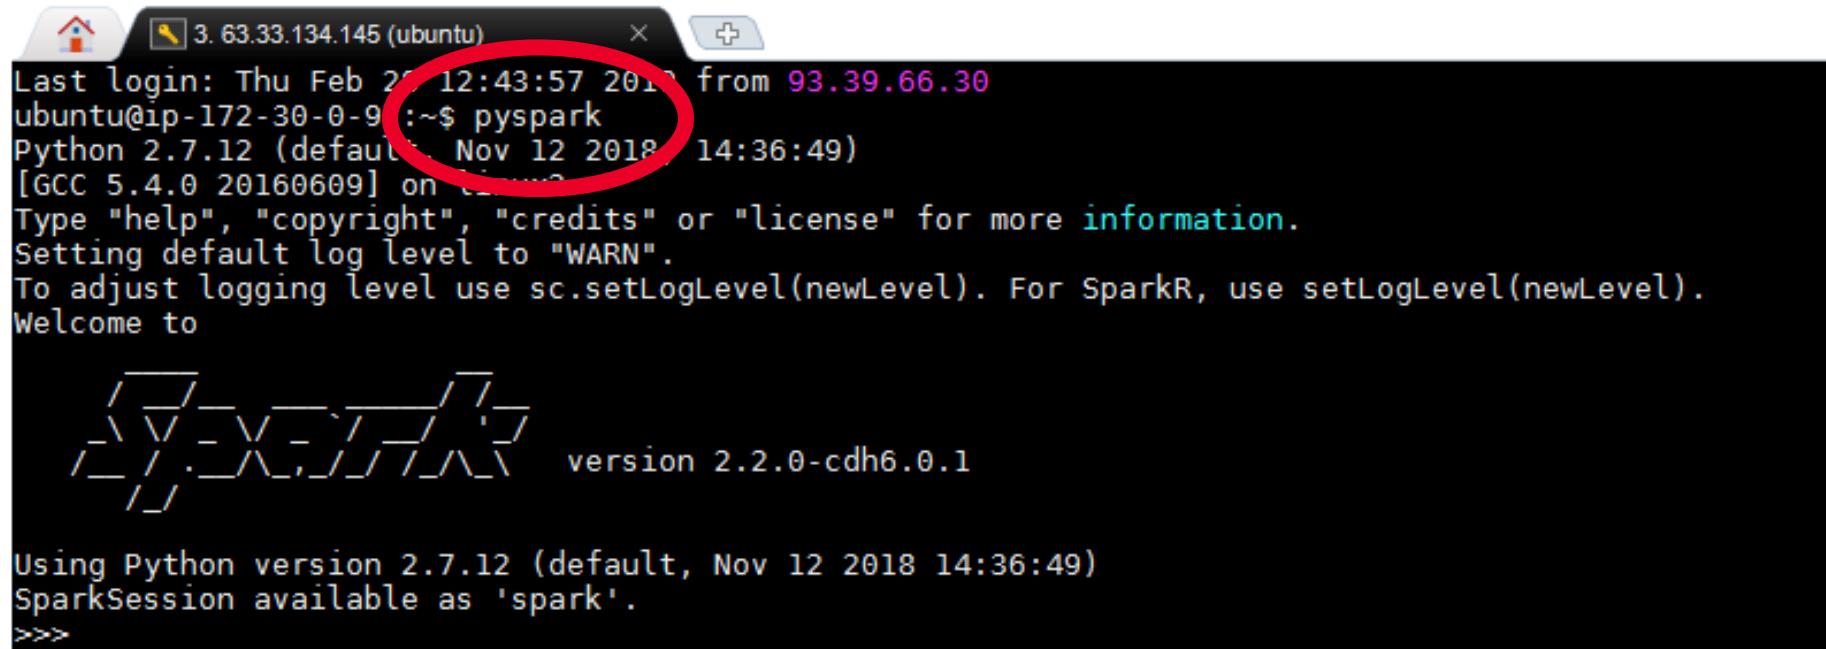
```
Last login: Fri Feb 22 08:34:07 2019 from 93.39.66.30
ubuntu@ip-172-30-0-97:~$ spark-shell
Setting default log level to "WARN"
To adjust logging level use setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Spark context Web UI available at http://172.30.0.97:4040
Spark context available as 'sc' (master = yarn, app id = application_1550659209716_0030).
Spark session available as 'spark'.
Welcome to

    \_____/ .-X-. / \ \    version 2.2.0-cdh6.0.1
   / \ \ .-X-. / \ \
   \ \ \ .-X-. / \ \
    \ \ \ .-X-. / \ \
     \ \ \ .-X-. / \ \
      \ \ \ .-X-. / \ \
       \ \ \ .-X-. / \ \
        \ \ \ .-X-. / \ \
         \ \ \ .-X-. / \ \
          \ \ \ .-X-. / \ \
           \ \ \ .-X-. / \ \
            \ \ \ .-X-. / \ \
             \ \ \ .-X-. / \ \
              \ \ \ .-X-. / \ \
               \ \ \ .-X-. / \ \
                \ \ \ .-X-. / \ \
                 \ \ \ .-X-. / \ \
                  \ \ \ .-X-. / \ \
                   \ \ \ .-X-. / \ \
                    \ \ \ .-X-. / \ \
                     \ \ \ .-X-. / \ \
                      \ \ \ .-X-. / \ \
                       \ \ \ .-X-. / \ \
                        \ \ \ .-X-. / \ \
                         \ \ \ .-X-. / \ \
                          \ \ \ .-X-. / \ \
                           \ \ \ .-X-. / \ \
                            \ \ \ .-X-. / \ \
                             \ \ \ .-X-. / \ \
                              \ \ \ .-X-. / \ \
                               \ \ \ .-X-. / \ \
                                \ \ \ .-X-. / \ \
                                 \ \ \ .-X-. / \ \
                                  \ \ \ .-X-. / \ \
                                   \ \ \ .-X-. / \ \
                                    \ \ \ .-X-. / \ \
                                     \ \ \ .-X-. / \ \
                                      \ \ \ .-X-. / \ \
                                       \ \ \ .-X-. / \ \
                                        \ \ \ .-X-. / \ \
                                         \ \ \ .-X-. / \ \
                                          \ \ \ .-X-. / \ \
                                           \ \ \ .-X-. / \ \
                                            \ \ \ .-X-. / \ \
                                             \ \ \ .-X-. / \ \
                                              \ \ \ .-X-. / \ \
                                               \ \ \ .-X-. / \ \
                                                \ \ \ .-X-. / \ \
                                                 \ \ \ .-X-. / \ \
                                                  \ \ \ .-X-. / \ \
                                                   \ \ \ .-X-. / \ \
                                                    \ \ \ .-X-. / \ \
                                                     \ \ \ .-X-. / \ \
                                                      \ \ \ .-X-. / \ \
                                                       \ \ \ .-X-. / \ \
                                                        \ \ \ .-X-. / \ \
                                                         \ \ \ .-X-. / \ \
                                                          \ \ \ .-X-. / \ \
                                                           \ \ \ .-X-. / \ \
                                                            \ \ \ .-X-. / \ \
                                                             \ \ \ .-X-. / \ \
                                                              \ \ \ .-X-. / \ \
                                                               \ \ \ .-X-. / \ \
                                                                \ \ \ .-X-. / \ \
                                                                 \ \ \ .-X-. / \ \
                                                                  \ \ \ .-X-. / \ \
                                                                   \ \ \ .-X-. / \ \
                                                                    \ \ \ .-X-. / \ \
                                                                     \ \ \ .-X-. / \ \
                                                                      \ \ \ .-X-. / \ \
                                                                       \ \ \ .-X-. / \ \
                                                                        \ \ \ .-X-. / \ \
                                                                         \ \ \ .-X-. / \ \
                                                                          \ \ \ .-X-. / \ \
                                                                           \ \ \ .-X-. / \ \
................................................................Using Scala version 2.11.8 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_141)
Type in expressions to have them evaluated.
Type :help for more information.

scala>
```

Come useremo Spark nel laboratorio di oggi?

Si può usare anche Python (pyspark)



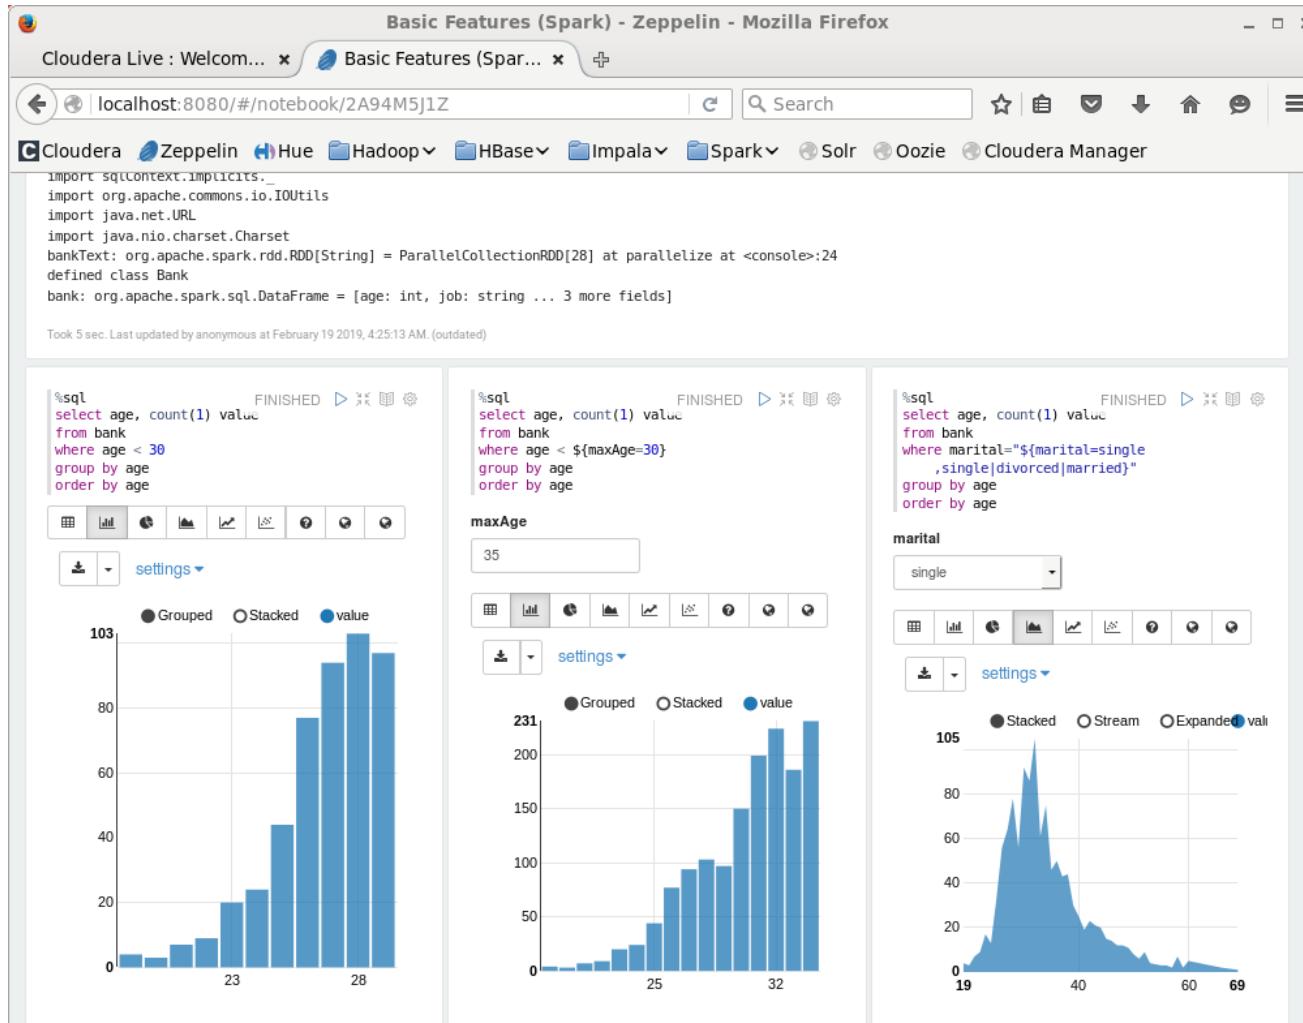
```
Last login: Thu Feb 27 12:43:57 2019 from 93.39.66.30
ubuntu@ip-172-30-0-9:~$ pyspark
Python 2.7.12 (default, Nov 12 2018 14:36:49)
[GCC 5.4.0 20160609] on linux2
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to

    \_____
   /       \
  /  _    _ \
 /  /\_/\_ \
/  /  \  \_ \
 \  \  /  \_ \
  \  /  /  \_ \
   \_ \_ \_ \_ \
                version 2.2.0-cdh6.0.1

Using Python version 2.7.12 (default, Nov 12 2018 14:36:49)
SparkSession available as 'spark'.
>>>
```

Come useremo Spark nel laboratorio di oggi?

Zeppelin Notebook



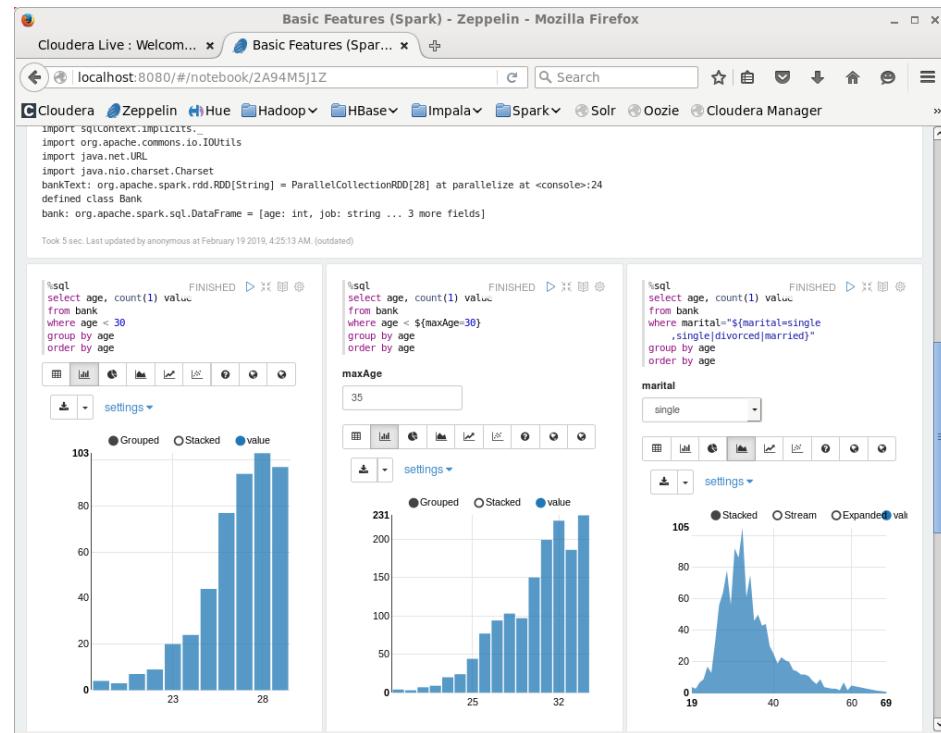
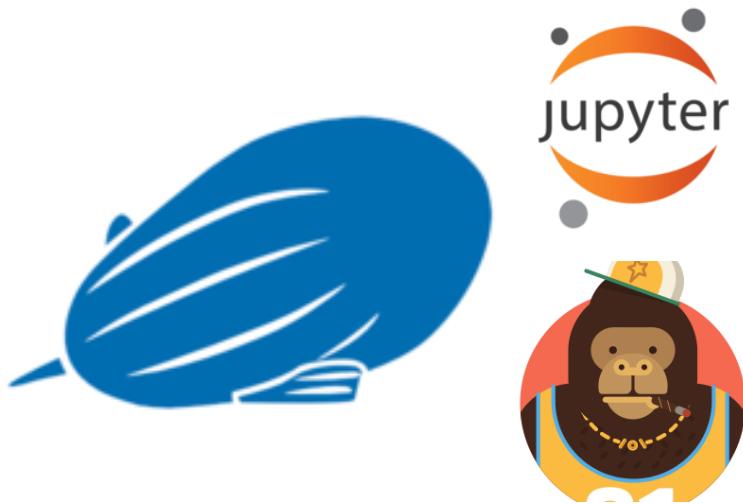
I Notebook

Un ambiente per il «literate programming»

Un'interfaccia che permette di mescolare frammenti eseguibili di programma con grafici, testo e altri contenuti.

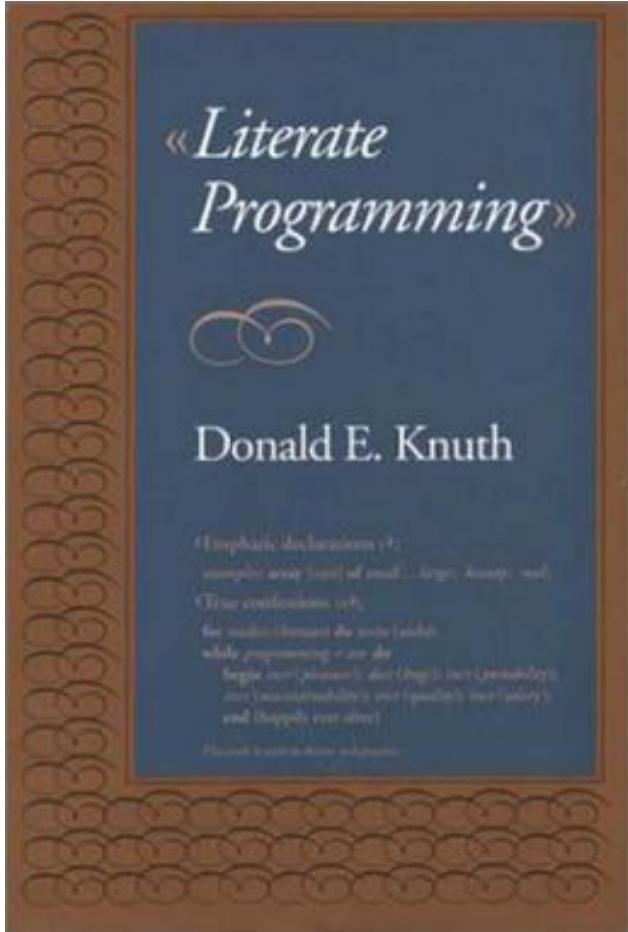
I notebook per Apache Spark sono:

- *Jupyter Notebook*
- *Apache Zeppelin*
- *Apache Spark Notebook*



Literate programming

Donald E. Knuth



Laboratorio

Obiettivi per questa mattina

- Ripasso di Scala
- Prendere familiarità con *Zeppelin Notebook*



Scala Cheat Sheet

variances

| | |
|---------------------------------------------|--------------------|
| Given $T <: P$ | |
| When C is covariant then $C[T] <: C[P]$ | |
| When C is contravariant then $C[T] >: C[P]$ | |
| class C[+T] | C is covariant |
| class C[-T] | C is contravariant |
| class C[T] | C is invariant |

variables and functions

| | |
|--------------------------------------------|----------------------------------------|
| val x = 1 | constant, evaluated immediately |
| var y = 2 | variable |
| def z = 2 | evaluated when called |
| lazy val w = 2 | evaluated once when needed |
| def f(x: Int) = { x*x } | define function |
| def f(x: R) | call by value |
| def f(x: => R) | call by name (lazy parameters) |
| (1 to 5).map(_*2) | anonymous function |
| | (positionally matched arguments) |
| (1 to 5).map(2*) | anonymous function |
| | (bound infix method) |
| (1 to 5).map(x => x*x) | anonymous function |
| | (to use an arg twice, have to name it) |
| def mapmake[T](g:T=>T)(seq: List[T]) = | generic type |
| seq.map(g) | |
| def sum(args: Int*) = args.reduceLeft(_+_) | varargs |

control constructs

| | |
|-----------------------------------------|-------------------|
| if (condition) this else that | if then else |
| while (x < 5) { println(x); x += 1} | while loop |
| do { println(x); x += 1 } while (x < 5) | do while loop |
| for (x <- xs if x%2 == 0) yield x*10 | for comprehension |
| for (i <- 1 to 5) { | for loop |
| println(i) | |
| } | |

object orientation

| | |
|----------------------------|------------------------------------|
| class C(x: R) | class with private param |
| class C(val x: R) | class with public param |
| new { ... } | anonymous class |
| abstract class D { ... } | abstract class |
| class C extends D { ... } | inheritance |
| class C(x: R) extends D(x) | inheritance and constructor params |
| object O extends D { ... } | singleton |
| trait T { ... } | interfaces with implementation |
| class C extends D with T | using multiple traits |

gutefrage.net

Zeppelin Notebook

Manuale online

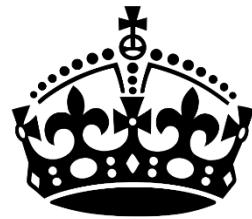
The screenshot shows two side-by-side views of the Apache Zeppelin website.

Left View: Install Page

- Header:** Zeppelin Quick Start Usage Setup Interpreter More
- Title:** Install
- Content:**
 - Requirements
 - Downloading Binary Package
 - Building Zeppelin from source
 - Starting Apache Zeppelin
 - Start Apache Zeppelin with a service manager
 - Next Steps
- Welcome Message:** Welcome to Apache Zeppelin! On this page are instructions to help you...
- Requirements Section:** Apache Zeppelin officially supports and is tested on the following environments:

| Name | Value |
|------------|-----------------------------------------------------------|
| Oracle JDK | 1.7 (set JAVA_HOME) |
| OS | Mac OSX Ubuntu 14.X CentOS 6.X Windows 7 Pro SP1 |

<https://zeppelin.apache.org/>



**KEEP
CALM
AND
START
CODING**